The markovchain Package: A Package for Easily Handling Discrete Markov Chains in R

Giorgio Alfredo Spedicato, Tae Seung Kang, Sai Bhargav Yalamanchi, Deepak Yadav, Ignacio Cordón

Abstract

The **markovchain** package aims to fill a gap within the R framework providing S4 classes and methods for easily handling discrete time Markov chains, homogeneous and simple inhomogeneous ones as well as continuous time Markov chains. The S4 classes for handling and analysing discrete and continuous time Markov chains are presented, as well as functions and method for performing probabilistic and statistical analysis. Finally, some examples in which the package's functions are applied to Economics, Finance and Natural Sciences topics are shown.

Keywords: discrete time Markov chains, continuous time Markov chains, transition matrices, communicating classes, periodicity, first passage time, stationary distributions.

1. Introduction

Markov chains represent a class of stochastic processes of great interest for the wide spectrum of practical applications. In particular, discrete time Markov chains (DTMC) permit to model the transition probabilities between discrete states by the aid of matrices. Various R packages deal with models that are based on Markov chains:

- **msm** (Jackson 2011) handles Multi-State Models for panel data.
- **mcmcR** (Geyer and Johnson 2013) implements Monte Carlo Markov Chain approach.
- **hmm** (Himmelmann and www.linhi.com 2010) fits hidden Markov models with covariates.
- **mstate** fits 'Multi-State Models based on Markov chains for survival analysis (de Wreede, Fiocco, and Putter 2011).

Nevertheless, the R statistical environment (R Core Team 2013) seems to lack a simple package that coherently defines S4 classes for discrete Markov chains and allows to perform probabilistic analysis, statistical inference and applications. For the sake of completeness, **markovchain** is the second package specifically dedicated to DTMC analysis, being **DTMCPack** (Nicholson 2013) the first one. Notwithstanding, **markovchain** package (Spedicato 2017) aims to offer more flexibility in handling DTMC than other existing solutions, providing S4 classes for both homogeneous and non-homogeneous Markov chains as well as methods suited to perform statistical and probabilistic analysis.

The **markovchain** package depends on the following R packages: **expm** (Goulet, Dutang, Maechler, Firth, Shapira, Stadelmann, and expm-developers@lists.R-forge.R-project.org 2013)

to perform efficient matrices powers; **igraph** (Csardi and Nepusz 2006) to perform pretty plotting of `markovchain` objects and **matlab** (Roebuck 2011), that contains functions for matrix management and calculations that emulate those within MATLAB environment. Moreover, other scientific softwares provide functions specifically designed to analyze DTMC, as Mathematica 9 (Wolfram Research 2013b).

The paper is structured as follows: Section 2 briefly reviews mathematics and definitions regarding DTMC, Section 3 discusses how to handle and manage Markov chain objects within the package, Section 4 and Section 5 show how to perform probabilistic and statistical modelling, while Section 6 presents some applied examples from various fields analyzed by means of the **markovchain** package.

2. Review of core mathematical concepts

2.1. General Definitions

A DTMC is a sequence of random variables $X_1, X_2, \ldots, X_n, \ldots$ characterized by the Markov property (also known as memoryless property, see Equation 1). The Markov property states that the distribution of the forthcoming state X_{n+1} depends only on the current state X_n and doesn't depend on the previous ones $X_{n-1}, X_{n-2}, \ldots, X_1$.

$$Pr\left(X_{n+1} = x_{n+1} \,|\, X_1 = x_1, X_2 = x_2, ..., X_n = x_n\right) = Pr\left(X_{n+1} = x_{n+1} \,|\, X_n = x_n\right). \quad (1)$$

The set of possible states $S = \{s_1, s_2, ..., s_r\}$ of X_n can be finite or countable and it is named the state space of the chain.

The chain moves from one state to another (this change is named either 'transition' or 'step') and the probability p_{ij} to move from state s_i to state s_j in one step is named transition probability:

$$p_{ij} = Pr\left(X_1 = s_j \,|\, X_0 = s_i\right). \quad (2)$$

The probability of moving from state i to j in n steps is denoted by $p_{ij}^{(n)} = Pr\left(X_n = s_j \,|\, X_0 = s_i\right)$. A DTMC is called time-homogeneous if the property shown in Equation 3 holds. Time homogeneity implies no change in the underlying transition probabilities as time goes on.

$$Pr\left(X_{n+1} = s_j \,|\, X_n = s_i\right) = Pr\left(X_n = s_j \,|\, X_{n-1} = s_i\right). \quad (3)$$

If the Markov chain is time-homogeneous, then $p_{ij} = Pr\left(X_{k+1} = s_j \,|\, X_k = s_i\right)$ and $p_{ij}^{(n)} = Pr\left(X_{n+k} = s_j \,|\, X_k = s_i\right)$, where $k > 0$.

The probability distribution of transitions from one state to another can be represented into a transition matrix $P = (p_{ij})_{i,j}$, where each element of position (i, j) represents the transition probability p_{ij}. E.g., if $r = 3$ the transition matrix P is shown in Equation 4

$$P = \begin{bmatrix} p_{11} & p_{12} & p_{13} \\ p_{21} & p_{22} & p_{23} \\ p_{31} & p_{32} & p_{33} \end{bmatrix}. \quad (4)$$

The distribution over the states can be written in the form of a stochastic row vector x (the term stochastic means that $\sum_i x_i = 1, x_i \geq 0$): e.g., if the current state of x is s_2, $x = (0\ 1\ 0)$. As a consequence, the relation between $x^{(1)}$ and $x^{(0)}$ is $x^{(1)} = x^{(0)}P$ and, recursively, we get $x^{(2)} = x^{(0)}P^2$ and $x^{(n)} = x^{(0)}P^n$, $n > 0$.

DTMC are explained in most theory books on stochastic processes, see Brémaud (1999) and Dobrow (2016) for example. Valuable references online available are: Konstantopoulos (2009), Snell (1999) and Bard (2000).

2.2. Properties and classification of states

A state s_j is said accessible from state s_i (written $s_i \rightarrow s_j$) if a system starting in state s_i has a positive probability to reach the state s_j at a certain point, i.e., $\exists n > 0 : p_{ij}^n > 0$. If both $s_i \rightarrow s_j$ and $s_j \rightarrow s_i$, then s_i and s_j are said to communicate.

A communicating class is defined to be a set of states that communicate. A DTMC can be composed by one or more communicating classes. If the DTMC is composed by only one communicating class (i.e., if all states in the chain communicate), then it is said irreducible. A communicating class is said to be closed if no states outside of the class can be reached from any state inside it.

If $p_{ii} = 1$, s_i is defined as absorbing state: an absorbing state corresponds to a closed communicating class composed by one state only.

The canonical form of a DTMC transition matrix is a matrix having a block form, where the closed communicating classes are shown at the beginning of the diagonal matrix.

A state s_i has period k_i if any return to state s_i must occur in multiplies of k_i steps, that is $k_i = gcd\{n : Pr(X_n = s_i | X_0 = s_i) > 0\}$, where gcd is the greatest common divisor. If $k_i = 1$ the state s_i is said to be aperiodic, else if $k_i > 1$ the state s_i is periodic with period k_i. Loosely speaking, s_i is periodic if it can only return to itself after a fixed number of transitions $k_i > 1$ (or multiple of k_i), else it is aperiodic.

If states s_i and s_j belong to the same communicating class, then they have the same period k_i. As a consequence, each of the states of an irreducible DTMC share the same periodicity. This periodicity is also considered the DTMC periodicity. It is possible to classify states according to their periodicity. Let $T^{x \rightarrow x}$ is the number of periods to go back to state x knowing that the chain starts in x.

- A state x is recurrent if $P(T^{x \rightarrow x} < +\infty) = 1$ (equivalently $P(T^{x \rightarrow x} = +\infty) = 0$). In addition:

    1. A state x is null recurrent if in addition $E(T^{x \rightarrow x}) = +\infty$.
    2. A state x is positive recurrent if in addition $E(T^{x \rightarrow x}) < +\infty$.
    3. A state x is absorbing if in addition $P(T^{x \rightarrow x} = 1) = 1$.

- A state x is transient if $P(T^{x \rightarrow x} < +\infty) < 1$ (equivalently $P(T^{x \rightarrow x} = +\infty) > 0$).

It is possible to analyze the timing to reach a certain state. The first passage time (or hitting time) from state s_i to state s_j is the number T_{ij} of steps taken by the chain until it arrives for the first time to state s_j, given that $X_0 = s_i$. The probability distribution of T_{ij} is defined

by Equation 5

$$h_{ij}^{(n)} = Pr\,(T_{ij} = n) = Pr\,(X_n = s_j, X_{n-1} \neq s_j, \ldots, X_1 \neq s_j | X_0 = s_i) \tag{5}$$

and can be found recursively using Equation 6, given that $h_{ij}^{(n)} = p_{ij}$.

$$h_{ij}^{(n)} = \sum_{k \in S - \{s_j\}} p_{ik} h_{kj}^{(n-1)}. \tag{6}$$

A commonly used quantity related to h is its average value, i.e. the *mean first passage time* (also expected hitting time), namely $\bar{h}_{ij} = \sum_{n=1\ldots\infty} n\, h_{ij}^{(n)}$.

If in the definition of the first passage time we let $s_i = s_j$, we obtain the first recurrence time $T_i = \inf\{n \geq 1 : X_n = s_i | X_0 = s_i\}$. We could also ask ourselves which is the *mean recurrence time*, an average of the mean first recurrence times:

$$r_i = \sum_{k=1}^{\infty} k \cdot P(T_i = k)$$

Revisiting the definition of recurrence and transience: a state s_i is said to be recurrent if it is visited infinitely often, i.e., $Pr(T_i < +\infty | X_0 = s_i) = 1$. On the opposite, s_i is called transient if there is a positive probability that the chain will never return to s_i, i.e., $Pr(T_i = +\infty | X_0 = s_i) > 0$.

Given a time homogeneous Markov chain with transition matrix P, a stationary distribution z is a stochastic row vector such that $z = z \cdot P$, where $0 \leq z_j \leq 1 \; \forall j$ and $\sum_j z_j = 1$.

If a DTMC $\{X_n\}$ is irreducible and aperiodic, then it has a limit distribution and this distribution is stationary. As a consequence, if P is the $k \times k$ transition matrix of the chain and $z = (z_1, ..., z_k)$ is the unique eigenvector of P such that $\sum_{i=1}^{k} z_i = 1$, then we get

$$\lim_{n \to \infty} P^n = Z, \tag{7}$$

where Z is the matrix having all rows equal to z. The stationary distribution of $\{X_n\}$ is represented by z.

A matrix A is called primitive if all of its entries are strictly positive, and we write it $A > 0$. If the transition matrix P for a DTMC has some primitive power, i.e. it exists $m > 0 : P^m > 0$, then the DTMC is said to be regular. In fact being regular is equivalent to being irreducible and aperiodic. All regular DTMCs are irreducible. The counterpart is not true.

Given two absorbing states s_A (source) and s_B (sink), the *committor probability* $q_j^{(AB)}$ is the probability that a process starting in state s_i is absorbed in state s_B (rather than s_A) (Noé, Schütte, Vanden-Eijnden, Reich, and Weikl 2009). It can be computed via

$$q_j^{(AB)} = \sum_{k \ni A,B} P_{jk} q_k^{(AB)} \quad \text{with} \quad q_A^{(AB)} = 0 \quad \text{and} \quad q_B^{(AB)} = 1 \tag{8}$$

Note we can also define the hitting probability from i to j as the probability of ever reaching the state j if our initial state is i:

$$h_{i,j} = Pr(T_{ij} < \infty) = \sum_{n=0}^{\infty} h_{ij}^{(n)} \qquad (9)$$

In a DTMC with finite set of states, we know that a transient state communicates at least with one recurrent state. If the chain starts in a transient element, once it hits a recurrent state, it is going to be caught in its recurrent state, and we cannot expect it would go back to the initial state. Given a transient state i we can define the *absorption probability* to the recurrent state j as the probability that the first recurrent state that the Markov chain visits (and therefore gets absorbed by its recurrent class) is j, f_i^*j. We can also define the *mean absorption time* as the mean number of steps the transient state i would take until it hits any recurrent state, b_i.

2.3. A short example

Consider the following numerical example. Suppose we have a DTMC with a set of 3 possible states $S = \{s_1, s_2, s_3\}$. Let the transition matrix be:

$$P = \begin{bmatrix} 0.5 & 0.2 & 0.3 \\ 0.15 & 0.45 & 0.4 \\ 0.25 & 0.35 & 0.4 \end{bmatrix}. \qquad (10)$$

In P, $p_{11} = 0.5$ is the probability that $X_1 = s_1$ given that we observed $X_0 = s_1$ is 0.5, and so on. It is easy to see that the chain is irreducible since all the states communicate (it is made by one communicating class only).

Suppose that the current state of the chain is $X_0 = s_2$, i.e., $x^{(0)} = (0\,1\,0)$, then the probability distribution of states after 1 and 2 steps can be computed as shown in Equations (11) and (12).

$$x^{(1)} = (0\,1\,0) \begin{bmatrix} 0.5 & 0.2 & 0.3 \\ 0.15 & 0.45 & 0.4 \\ 0.25 & 0.35 & 0.4 \end{bmatrix} = (0.15\,0.45\,0.4). \qquad (11)$$

$$x^{(n)} = x^{(n-1)}P \rightarrow (0.15\,0.45\,0.4) \begin{bmatrix} 0.5 & 0.2 & 0.3 \\ 0.15 & 0.45 & 0.4 \\ 0.25 & 0.35 & 0.4 \end{bmatrix} = (0.2425\,0.3725\,0.385). \qquad (12)$$

If we were interested in the probability of being in the state s_3 in the second step, then $Pr(X_2 = s_3 \,|\, X_0 = s_2) = 0.385$.

3. The structure of the package

3.1. Creating markovchain objects

The package is loaded within the R command line as follows:

```
R> library("markovchain")
```

```
Loading required package: matlab
```

```
Attaching package: 'matlab'
```

```
The following object is masked from 'package:stats':

    reshape
```

```
The following objects are masked from 'package:utils':

    find, fix
```

```
The following object is masked from 'package:base':

    sum
```

The `markovchain` and `markovchainList` S4 classes (Chambers 2008) are defined within the **markovchain** package as displayed:

```
Class "markovchain" [package "markovchain"]
```

```
Slots:
```

```
Name:           states              byrow transitionMatrix             name
Class:       character            logical             matrix        character
```

```
Class "markovchainList" [package "markovchain"]
```

```
Slots:
```

```
Name:  markovchains           name
Class:         list      character
```

The first class has been designed to handle homogeneous Markov chain processes, while the latter (which is itself a list of `markovchain` objects) has been designed to handle non-homogeneous Markov chains processes.

Any element of `markovchain` class is comprised by following slots:

1. `states`: a character vector, listing the states for which transition probabilities are defined.
2. `byrow`: a logical element, indicating whether transition probabilities are shown by row or by column.
3. `transitionMatrix`: the probabilities of the transition matrix.
4. `name`: optional character element to name the DTMC.

The `markovchainList` objects are defined by following slots:

1. `markovchains`: a list of `markovchain` objects.
2. `name`: optional character element to name the DTMC.

The `markovchain` objects can be created either in a long way, as the following code shows

```
R> weatherStates <- c("sunny", "cloudy", "rain")
R> byRow <- TRUE
R> weatherMatrix <- matrix(data = c(0.70, 0.2, 0.1,
+                     0.3, 0.4, 0.3,
+                     0.2, 0.45, 0.35), byrow = byRow, nrow = 3,
+                   dimnames = list(weatherStates, weatherStates))
R> mcWeather <- new("markovchain", states = weatherStates, byrow = byRow,
+               transitionMatrix = weatherMatrix, name = "Weather")
```

or in a shorter way, displayed below

```
R> mcWeather <- new("markovchain", states = c("sunny", "cloudy", "rain"),
+               transitionMatrix = matrix(data = c(0.70, 0.2, 0.1,
+                     0.3, 0.4, 0.3,
+                     0.2, 0.45, 0.35), byrow = byRow, nrow = 3),
+                 name = "Weather")
```

When `new("markovchain")` is called alone, a default Markov chain is created.

```
R> defaultMc <- new("markovchain")
```

The quicker way to create `markovchain` objects is made possible thanks to the implemented `initialize` S4 method that checks that:

- the `transitionMatrix` to be a transition matrix, i.e., all entries to be probabilities and either all rows or all columns to sum up to one.
- the columns and rows names of `transitionMatrix` to be defined and to coincide with `states` vector slot.

The `markovchain` objects can be collected in a list within `markovchainList` S4 objects as following example shows.

```
R> mcList <- new("markovchainList", markovchains = list(mcWeather, defaultMc),
+                 name = "A list of Markov chains")
```

3.2. Handling markovchain objects

Table 1 lists which methods handle and manipulate `markovchain` objects.

Method	Purpose
*	Direct multiplication for transition matrices.
^	Compute the power `markovchain` of a given one.
[	Direct access to the elements of the transition matrix.
==	Equality operator between two transition matrices.
!=	Inequality operator between two transition matrices.
as	Operator to convert `markovchain` objects into `data.frame` and `table` object.
dim	Dimension of the transition matrix.
names	Equal to `states`.
names<-	Change the `states` name.
name	Get the name of `markovchain object`.
name<-	Change the name of `markovchain object`.
plot	`plot` method for `markovchain` objects.
print	`print` method for `markovchain` objects.
show	`show` method for `markovchain` objects.
sort	`sort` method for `markovchain` objects, in terms of their states.
states	Name of the transition states.
t	Transposition operator (which switches `byrow` 'slot value and modifies the transition matrix coherently).

Table 1: **markovchain** methods for handling `markovchain` objects.

The examples that follow shows how operations on `markovchain` objects can be easily performed. For example, using the previously defined matrix we can find what is the probability distribution of expected weather states in two and seven days, given the actual state to be cloudy.

```
R> initialState <- c(0, 1, 0)
R> after2Days <- initialState * (mcWeather * mcWeather)
R> after7Days <- initialState * (mcWeather ^ 7)
R> after2Days

    sunny cloudy  rain
[1,]  0.39  0.355 0.255


R> round(after7Days, 3)

    sunny cloudy  rain
[1,] 0.462  0.319 0.219
```

A similar answer could have been obtained defining the vector of probabilities as a column vector. A column - defined probability matrix could be set up either creating a new matrix or transposing an existing `markovchain` object thanks to the `t` method.

```
R> initialState <- c(0, 1, 0)
R> after2Days <- (t(mcWeather) * t(mcWeather)) * initialState
R> after7Days <- (t(mcWeather) ^ 7) * initialState
R> after2Days

        [,1]
sunny  0.390
cloudy 0.355
rain   0.255

R> round(after7Days, 3)

        [,1]
sunny  0.462
cloudy 0.319
rain   0.219
```

The initial state vector previously shown can not necessarily be a probability vector, as the code that follows shows:

```
R> fvals<-function(mchain,initialstate,n) {
+   out<-data.frame()
+   names(initialstate)<-names(mchain)
+   for (i in 0:n)
+   {
+     iteration<-initialstate*mchain^(i)
+     out<-rbind(out,iteration)
+   }
+   out<-cbind(out, i=seq(0,n))
+   out<-out[,c(4,1:3)]
+   return(out)
+ }
R> fvals(mchain=mcWeather,initialstate=c(90,5,5),n=4)

  i     sunny    cloudy      rain
1 0 90.00000   5.00000   5.00000
2 1 65.50000  22.25000  12.25000
3 2 54.97500  27.51250  17.51250
4 3 50.23875  29.88063  19.88062
5 4 48.10744  30.94628  20.94628
```

Basic methods have been defined for `markovchain` objects to quickly get states and transition matrix dimension.

```
R> states(mcWeather)

[1] "sunny"  "cloudy" "rain"

R> names(mcWeather)

[1] "sunny"  "cloudy" "rain"

R> dim(mcWeather)

[1] 3
```

Methods are available to set and get the name of `markovchain` object.

```
R> name(mcWeather)

[1] "Weather"

R> name(mcWeather) <- "New Name"
R> name(mcWeather)

[1] "New Name"
```

Also it is possible to alphabetically sort the transition matrix:

```
R> markovchain:::sort(mcWeather)

New Name
 A  3 - dimensional discrete Markov Chain defined by the following states:
 cloudy, rain, sunny
 The transition matrix  (by rows)  is defined as follows:
      cloudy rain sunny
cloudy   0.40 0.30   0.3
rain     0.45 0.35   0.2
sunny    0.20 0.10   0.7
```

A direct access to transition probabilities is provided both by `transitionProbability` method and `"["` method.

```
R> transitionProbability(mcWeather, "cloudy", "rain")

[1] 0.3

R> mcWeather[2,3]

[1] 0.3
```

The transition matrix of a `markovchain` object can be displayed using `print` or `show` methods (the latter being less verbose). Similarly, the underlying transition probability diagram can be plotted by the use of `plot` method (as shown in Figure 1) which is based on **igraph** package (Csardi and Nepusz 2006). `plot` method for `markovchain` objects is a wrapper of `plot.igraph` for `igraph` S4 objects defined within the **igraph** package. Additional parameters can be passed to `plot` function to control the network graph layout. There are also **diagram** and **DiagrammeR** ways available for plotting as shown in Figure 2. The `plot` function also uses `communicatingClasses` function to separate out states of different communicating classes. All states that belong to one class have same color.

```
R> print(mcWeather)

        sunny cloudy rain
sunny    0.7   0.20 0.10
cloudy   0.3   0.40 0.30
rain     0.2   0.45 0.35

R> show(mcWeather)

New Name
 A  3 - dimensional discrete Markov Chain defined by the following states:
 sunny, cloudy, rain
 The transition matrix  (by rows)  is defined as follows:
        sunny cloudy rain
sunny    0.7   0.20 0.10
cloudy   0.3   0.40 0.30
rain     0.2   0.45 0.35


Attaching package: 'igraph'

The following objects are masked from 'package:stats':

    decompose, spectrum

The following object is masked from 'package:base':

    union


Loading required package: shape
```

Import and export from some specific classes is possible, as shown in Figure 3 and in the following code.

```
R> mcDf <- as(mcWeather, "data.frame")
R> mcNew <- as(mcDf, "markovchain")
R> mcDf
```

Figure 1: Weather example. Markov chain plot

Figure 2: Weather example. Markov chain plot with diagram

```
      t0      t1 prob
1  sunny   sunny 0.70
2  sunny  cloudy 0.20
3  sunny    rain 0.10
4 cloudy   sunny 0.30
5 cloudy  cloudy 0.40
6 cloudy    rain 0.30
7   rain   sunny 0.20
8   rain  cloudy 0.45
9   rain    rain 0.35
```

```
R> mcIgraph <- as(mcWeather, "igraph")
```

```
R> if (requireNamespace("msm", quietly = TRUE)) {
+ require(msm)
+ Q <- rbind ( c(0, 0.25, 0, 0.25),
+             c(0.166, 0, 0.166, 0.166),
+             c(0, 0.25, 0, 0.25),
+             c(0, 0, 0, 0) )
+ cavmsm <- msm(state ~ years, subject = PTNUM, data = cav, qmatrix = Q, death = 4)
+ msmMc <- as(cavmsm, "markovchain")
+ msmMc
+   } else {
+   message("msm unavailable")
+   }
```

```
Loading required package: msm
```

```
Unnamed Markov chain
 A  4 - dimensional discrete Markov Chain defined by the following states:
 State 1, State 2, State 3, State 4
 The transition matrix  (by rows)  is defined as follows:
            State 1     State 2     State 3     State 4
State 1 0.853958721 0.08836953 0.01475543 0.04291632
State 2 0.155576908 0.56663284 0.20599563 0.07179462
State 3 0.009903994 0.07853691 0.65965727 0.25190183
State 4 0.000000000 0.00000000 0.00000000 1.00000000
```

from etm (now archived as of September 2020):

```
R> if (requireNamespace("etm", quietly = TRUE)) {
+ library(etm)
+ data(sir.cont)
+ sir.cont <- sir.cont[order(sir.cont$id, sir.cont$time), ]
+ for (i in 2:nrow(sir.cont)) {
+   if (sir.cont$id[i]==sir.cont$id[i-1]) {
```

```
+     if (sir.cont$time[i]==sir.cont$time[i-1]) {
+       sir.cont$time[i-1] <- sir.cont$time[i-1] - 0.5
+     }
+   }
+ }
+ tra <- matrix(ncol=3,nrow=3,FALSE)
+ tra[1, 2:3] <- TRUE
+ tra[2, c(1, 3)] <- TRUE
+ tr.prob <- etm::etm(sir.cont, c("0", "1", "2"), tra, "cens", 1)
+ tr.prob
+ etm2mc<-as(tr.prob, "markovchain")
+ etm2mc
+   } else {
+   message("etm unavailable")
+ }


Unnamed Markov chain
 A  3 - dimensional discrete Markov Chain defined by the following states:
 0, 1, 2
 The transition matrix  (by rows)  is defined as follows:
          0         1         2
0 0.0000000 0.5000000 0.5000000
1 0.5000000 0.0000000 0.5000000
2 0.3333333 0.3333333 0.3333333
```

Coerce from `matrix` method, as the code below shows, represents another approach to create a `markovchain` method starting from a given squared probability matrix.

```
R> myMatr<-matrix(c(.1,.8,.1,.2,.6,.2,.3,.4,.3), byrow=TRUE, ncol=3)
R> myMc<-as(myMatr, "markovchain")
R> myMc


Unnamed Markov chain
 A  3 - dimensional discrete Markov Chain defined by the following states:
 s1, s2, s3
 The transition matrix  (by rows)  is defined as follows:
    s1  s2  s3
s1 0.1 0.8 0.1
s2 0.2 0.6 0.2
s3 0.3 0.4 0.3
```

Non-homogeneous Markov chains can be created with the aid of `markovchainList` object. The example that follows arises from health insurance, where the costs associated to patients in a Continuous Care Health Community (CCHC) are modeled by a non-homogeneous Markov Chain, since the transition probabilities change by year. Methods explicitly written for `markovchainList` objects are: `print`, `show`, `dim` and `[`.
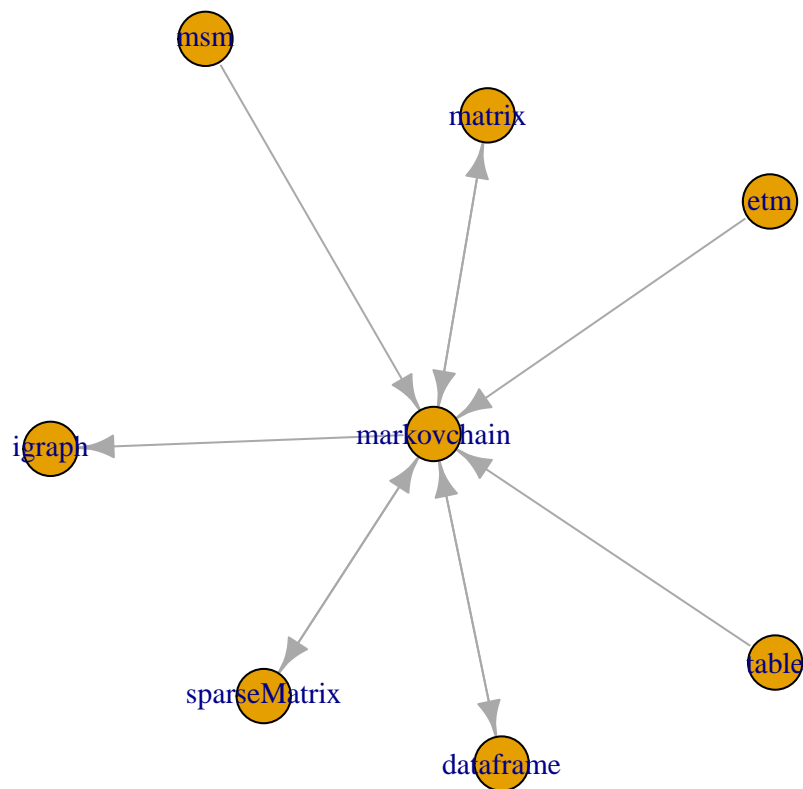
Import – Export from and to markovchain objects

Figure 3: The markovchain methods for import and export

```
R> stateNames = c("H", "I", "D")
R> Q0 <- new("markovchain", states = stateNames,
+         transitionMatrix =matrix(c(0.7, 0.2, 0.1,0.1, 0.6, 0.3,0, 0, 1),
+         byrow = TRUE, nrow = 3), name = "state t0")
R> Q1 <- new("markovchain", states = stateNames,
+         transitionMatrix = matrix(c(0.5, 0.3, 0.2,0, 0.4, 0.6,0, 0, 1),
+         byrow = TRUE, nrow = 3), name = "state t1")
R> Q2 <- new("markovchain", states = stateNames,
+         transitionMatrix = matrix(c(0.3, 0.2, 0.5,0, 0.2, 0.8,0, 0, 1),
+         byrow = TRUE,nrow = 3), name = "state t2")
R> Q3 <- new("markovchain", states = stateNames,
+          transitionMatrix = matrix(c(0, 0, 1, 0, 0, 1, 0, 0, 1),
+         byrow = TRUE, nrow = 3), name = "state t3")
R> mcCCRC <- new("markovchainList",markovchains = list(Q0,Q1,Q2,Q3),
+       name = "Continuous Care Health Community")
R> print(mcCCRC)

Continuous Care Health Community  list of Markov chain(s)
Markovchain  1
state t0
 A  3 - dimensional discrete Markov Chain defined by the following states:
 H, I, D
 The transition matrix  (by rows)  is defined as follows:
    H   I   D
H 0.7 0.2 0.1
I 0.1 0.6 0.3
D 0.0 0.0 1.0


Markovchain  2
state t1
 A  3 - dimensional discrete Markov Chain defined by the following states:
 H, I, D
 The transition matrix  (by rows)  is defined as follows:
    H   I   D
H 0.5 0.3 0.2
I 0.0 0.4 0.6
D 0.0 0.0 1.0


Markovchain  3
state t2
 A  3 - dimensional discrete Markov Chain defined by the following states:
 H, I, D
 The transition matrix  (by rows)  is defined as follows:
    H   I   D
H 0.3 0.2 0.5
I 0.0 0.2 0.8
D 0.0 0.0 1.0
```

```
Markovchain  4
state t3
 A  3 - dimensional discrete Markov Chain defined by the following states:
 H, I, D
 The transition matrix  (by rows)  is defined as follows:
  H I D
H 0 0 1
I 0 0 1
D 0 0 1
```

It is possible to perform direct access to `markovchainList` elements, as well as to determine the number of `markovchain` objects by which a `markovchainList` object is composed.

```
R> mcCCRC[[1]]
```

```
state t0
 A  3 - dimensional discrete Markov Chain defined by the following states:
 H, I, D
 The transition matrix  (by rows)  is defined as follows:
    H   I   D
H 0.7 0.2 0.1
I 0.1 0.6 0.3
D 0.0 0.0 1.0
```

```
R> dim(mcCCRC)
```

```
[1] 4
```

The `markovchain` package contains some data found in the literature related to DTMC models (see Section 6. Table 2 lists datasets and tables included within the current release of the package.

Dataset	Description
blanden	Mobility across income quartiles, Jo Blanden and Machin (2005).
craigsendi	CD4 cells, B. A. Craig and A. A. Sendi (2002).
kullback	raw transition matrices for testing homogeneity, Kullback, Kupperman, and Ku (1962).
preproglucacon	Preproglucacon DNA basis, P. J. Avery and D. A. Henderson (1999).
rain	Alofi Island rains, P. J. Avery and D. A. Henderson (1999).
holson	Individual states trajectories.
sales	Sales of six beverages in Hong Kong Ching, Ng, and Fung (2008).

Table 2: The **markovchain** `data.frame` and `table`.

Finally, Table 3 lists the demos included in the demo directory of the package.

R Code File	Description
`bard.R`	Structural analysis of Markov chains from Bard PPT.
`examples.R`	Notable Markov chains, e.g., The Gambler Ruin chain.
`quickStart.R`	Generic examples.
`extractMatrices.R`	Generic examples.

Table 3: The **markovchain** demos.

4. Probability with markovchain objects

The **markovchain** package contains functions to analyse DTMC from a probabilistic perspective. For example, the package provides methods to find stationary distributions and identifying absorbing and transient states. Many of these methods come from MATLAB listings that have been ported into R. For a full description of the underlying theory and algorithm the interested reader can overview the original MATLAB listings, Feres (2007) and Montgomery (2009).

Table 4 shows methods that can be applied on `markovchain` objects to perform probabilistic analysis.

Method	Returns
`absorbingStates`	the absorbing states of the transition matrix, if any.
`steadyStates`	the vector(s) of steady state(s) in matrix form.
`meanFirstPassageTime`	matrix or vector of mean first passage times.
`meanRecurrenceTime`	vector of mean number of steps to return to each recurrent state
`hittingProbabilities`	matrix of hitting probabilities for a Markov chain.
`meanAbsorptionTime`	expected number of steps for a transient state to be absorbed by any recurrent class
`absorptionProbabilities`	probabilities of transient states of being absorbed by each recurrent state
`committorAB`	committor probabilities
`communicatingClasses`	list of communicating classes. s_j, given actual state s_i.
`canonicForm`	the transition matrix into canonic form.
`is.accessible`	checks whether a state j is reachable from state i.
`is.irreducible`	checks whether a DTMC is irreducible.
`is.regular`	checks whether a DTMC is regular.
`period`	the period of an irreducible DTMC.
`recurrentClasses`	list of recurrent communicating classes.
`transientClasses`	list of transient communicating classes.
`recurrentStates`	the recurrent states of the transition matrix.
`transientStates`	the transient states of the transition matrix, if any.
`summary`	DTMC summary.

Table 4: **markovchain** methods: statistical operations.

4.1. Conditional distributions

The conditional distribution of weather states, given that current day's weather is sunny, is given by following code.

```
R> conditionalDistribution(mcWeather, "sunny")

 sunny cloudy   rain
   0.7    0.2    0.1
```

4.2. Stationary states

A stationary (steady state, or equilibrium) vector is a probability vector such that Equation 13 holds

$$
\begin{aligned}
0 \leq \pi_j \leq 1 \\
\sum_{j \in S} \pi_j = 1 \\
\pi \cdot P = \pi
\end{aligned}
\tag{13}
$$

Steady states are associated to P eigenvalues equal to one. We could be tempted to compute them solving the eigen values / vectors of the matrix and taking real parts (since if $u + iv$ is a eigen vector, for the matrix P, then $Re(u + iv) = u$ and $Im(u + iv) = v$ are eigen vectors) and normalizing by the vector sum, this carries some concerns:

1. If $u, v \in \mathbb{R}^n$ are linearly independent eigen vectors associated to 1 eigen value, $u + iv$, $u + iu$ are also linearly independent eigen vectors, and their real parts coincide. Clearly if we took real parts, we would be loosing an eigen vector, because we cannot know in advance if the underlying algorithm to compute the eigen vectors is going to output something similar to what we described. We should be agnostic to the underlying eigen vector computation algorithm.

2. Imagine the identity P of dimensions 2×2. Its eigen vectors associated to the 1 eigen value are $u = (1, 0)$ and $v = (0, 1)$. However, the underlying algorithm to compute eigen vectors could return $(1, -2)$ and $(-2, 1)$ instead, that are linear combinations of the aforementioned ones, and therefore eigen vectors. Normalizing by their sum, we would get: $(-1, 2)$ and $(2, -1)$, which obviously are not probability measures. Again, we should be agnostic to the underlying eigen computation algorithm.

3. Algorithms to compute eigen values / vectors are computationally expensive: they are iterative, and we cannot predict a fixed number of iterations for them. Moreover, each iteration takes $\mathcal{O}(m^2)$ or $\mathcal{O}(m^3)$ algorithmic complexity, with m the number of states.

We are going to use that every irreducible DTMC has a unique steady state, that is, if M is the matrix for an irreducible DTMC (all states communicate with each other), then it exists a unique $v \in \mathbb{R}^m$ such that:

$$
v \cdot M = v, \qquad \sum_{i=1}^{m} v_i = 1
$$

Also, we'll use that a steady state for a DTMC assigns 0 to the transient states. The canonical form of a (by row) stochastic matrix looks alike:

$$
\begin{pmatrix}
\begin{array}{c|c|c|c|c}
M_1 & 0 & 0 & \dots & 0 \\
\hline
0 & M_2 & 0 & \dots & 0 \\
\hline
0 & 0 & M_3 & \dots & 0 \\
\hline
\vdots & \vdots & \vdots & \ddots & \vdots \\
\hline
A_1 & A_2 & A_3 & \dots & R
\end{array}
\end{pmatrix}
$$

where M_i corresponds to irreducible sub-chains, the blocks A_i correspond to the transitions from transient states to each of the recurrent classes and R are the transitions from the transient states to themselves.

Also, we should note that a Markov chain has exactly the same name of steady states as recurrent classes. Therefore, we have coded the following algorithm [1]:

1. Identify the recurrent classes $[C_1, \dots, C_l]$ with `recurrentClasses` function.
2. Take each class C_i, compute the sub-matrix corresponding to it M_i.
3. Solve the system $v \cdot C_i = v$, $\sum_{j=1}^{|C_i|} v_j = 1$ which has a unique solution, for each $i = 1, \dots, l$.
4. Map each state v_i to the original order in P and assign a 0 to the slots corresponding to transient states in the matrix.

The result is returned in matrix form.

```
R> steadyStates(mcWeather)


        sunny    cloudy      rain
[1,] 0.4636364 0.3181818 0.2181818
```

It is possible for a Markov chain to have more than one stationary distribution, as the gambler ruin example shows.

```
R> gamblerRuinMarkovChain <- function(moneyMax, prob = 0.5) {
+   m <- matlab::zeros(moneyMax + 1)
+   m[1,1] <- m[moneyMax + 1,moneyMax + 1] <- 1
+   states <- as.character(0:moneyMax)
+   rownames(m) <- colnames(m) <- states
+
+   for(i in 2:moneyMax){
+     m[i,i-1] <- 1 - prob
+     m[i, i + 1] <- prob
+   }
+
```

---

[1] We would like to thank Prof. Christophe Dutang for his contributions to the development of this method. He coded a first improvement of the original `steadyStates` method and we could not have reached the current correctness without his previous work

```
+   new("markovchain", transitionMatrix = m,
+       name = paste("Gambler ruin", moneyMax, "dim", sep = " "))
+ }
R>
R> mcGR4 <- gamblerRuinMarkovChain(moneyMax = 4, prob = 0.5)
R> steadyStates(mcGR4)


     0 1 2 3 4
[1,] 0 0 0 0 1
[2,] 1 0 0 0 0
```

4.3. Classification of states

Absorbing states are determined by means of `absorbingStates` method.

```
R> absorbingStates(mcGR4)


[1] "0" "4"


R> absorbingStates(mcWeather)


character(0)
```

The key function in methods which need knowledge about communicating classes, recurrent states, transient states, is `.commclassKernel`, which is a modification of Tarjan's algorithm from Tarjan (1972). This `.commclassKernel` method gets a transition matrix of dimension n and returns a list of two items:

1. `classes`, an matrix whose (i, j) entry is `true` if s_i and s_j are in the same communicating class.
2. `closed`, a vector whose i-th entry indicates whether the communicating class to which i belongs is closed.

These functions are used by two other internal functions on which the `summary` method for `markovchain` objects works.

The example matrix used in Feres (2007) well exemplifies the purpose of the function.

```
R> P <- matlab::zeros(10)
R> P[1, c(1, 3)] <- 1/2;
R> P[2, 2] <- 1/3; P[2,7] <- 2/3;
R> P[3, 1] <- 1;
R> P[4, 5] <- 1;
R> P[5, c(4, 5, 9)] <- 1/3;
R> P[6, 6] <- 1;
R> P[7, 7] <- 1/4; P[7,9] <- 3/4;
```

```
R> P[8, c(3, 4, 8, 10)] <- 1/4;
R> P[9, 2] <- 1;
R> P[10, c(2, 5, 10)] <- 1/3;
R> rownames(P) <- letters[1:10]
R> colnames(P) <- letters[1:10]
R> probMc <- new("markovchain", transitionMatrix = P,
+                name = "Probability MC")
R> summary(probMc)


Probability MC  Markov chain that is composed by:
Closed classes:
a c
b g i
f
Recurrent classes:
{a,c},{b,g,i},{f}
Transient classes:
{d,e},{h},{j}
The Markov chain is not irreducible
The absorbing states are: f
```

All states that pertain to a transient class are named "transient" and a specific method has been written to elicit them.

```
R> transientStates(probMc)


[1] "d" "e" "h" "j"
```

`canonicForm` method that turns a Markov chain into its canonic form, reordering the states to have first the recurrent classes and then the transient states.

```
R> probMcCanonic <- canonicForm(probMc)
R> probMc


Probability MC
 A  10 - dimensional discrete Markov Chain defined by the following states:
 a, b, c, d, e, f, g, h, i, j
 The transition matrix  (by rows)  is defined as follows:
    a         b    c         d         e f         g    h         i         j
a 0.5 0.0000000 0.50 0.0000000 0.0000000 0 0.0000000 0.00 0.0000000 0.0000000
b 0.0 0.3333333 0.00 0.0000000 0.0000000 0 0.6666667 0.00 0.0000000 0.0000000
c 1.0 0.0000000 0.00 0.0000000 0.0000000 0 0.0000000 0.00 0.0000000 0.0000000
d 0.0 0.0000000 0.00 0.0000000 1.0000000 0 0.0000000 0.00 0.0000000 0.0000000
e 0.0 0.0000000 0.00 0.3333333 0.3333333 0 0.0000000 0.00 0.3333333 0.0000000
f 0.0 0.0000000 0.00 0.0000000 0.0000000 1 0.0000000 0.00 0.0000000 0.0000000
g 0.0 0.0000000 0.00 0.0000000 0.0000000 0 0.2500000 0.00 0.7500000 0.0000000
```

```
h 0.0 0.0000000 0.25 0.2500000 0.0000000 0 0.0000000 0.25 0.0000000 0.2500000
i 0.0 1.0000000 0.00 0.0000000 0.0000000 0 0.0000000 0.00 0.0000000 0.0000000
j 0.0 0.3333333 0.00 0.0000000 0.3333333 0 0.0000000 0.00 0.0000000 0.3333333
```

```
R> probMcCanonic
```

```
Probability MC
 A  10 - dimensional discrete Markov Chain defined by the following states:
 a, c, b, g, i, f, d, e, h, j
 The transition matrix  (by rows)  is defined as follows:
     a    c         b         g         i f         d         e    h         j
a 0.5 0.50 0.0000000 0.0000000 0.0000000 0 0.0000000 0.0000000 0.00 0.0000000
c 1.0 0.00 0.0000000 0.0000000 0.0000000 0 0.0000000 0.0000000 0.00 0.0000000
b 0.0 0.00 0.3333333 0.6666667 0.0000000 0 0.0000000 0.0000000 0.00 0.0000000
g 0.0 0.00 0.0000000 0.2500000 0.7500000 0 0.0000000 0.0000000 0.00 0.0000000
i 0.0 0.00 1.0000000 0.0000000 0.0000000 0 0.0000000 0.0000000 0.00 0.0000000
f 0.0 0.00 0.0000000 0.0000000 0.0000000 1 0.0000000 0.0000000 0.00 0.0000000
d 0.0 0.00 0.0000000 0.0000000 0.0000000 0 0.0000000 1.0000000 0.00 0.0000000
e 0.0 0.00 0.0000000 0.0000000 0.3333333 0 0.3333333 0.3333333 0.00 0.0000000
h 0.0 0.25 0.0000000 0.0000000 0.0000000 0 0.2500000 0.0000000 0.25 0.2500000
j 0.0 0.00 0.3333333 0.0000000 0.0000000 0 0.0000000 0.3333333 0.00 0.3333333
```

The function `is.accessible` permits to investigate whether a state s_j is accessible from state s_i, that is whether the probability to eventually reach s_j starting from s_i is greater than zero.

```
R> is.accessible(object = probMc, from = "a", to = "c")
```

```
[1] TRUE
```

```
R> is.accessible(object = probMc, from = "g", to = "c")
```

```
[1] FALSE
```

In Section 2.2 we observed that, if a DTMC is irreducible, all its states share the same periodicity. Then, the `period` function returns the periodicity of the DTMC, provided that it is irreducible. The example that follows shows how to find if a DTMC is reducible or irreducible by means of the function `is.irreducible` and, in the latter case, the method `period` is used to compute the periodicity of the chain.

```
R> E <- matrix(0, nrow = 4, ncol = 4)
R> E[1, 2] <- 1
R> E[2, 1] <- 1/3; E[2, 3] <- 2/3
R> E[3,2] <- 1/4; E[3, 4] <- 3/4
R> E[4, 3] <- 1
R>
R> mcE <- new("markovchain", states = c("a", "b", "c", "d"),
+        transitionMatrix = E,
+        name = "E")
R> is.irreducible(mcE)
```

Figure 4: Mathematica 9 example. Markov chain plot.

```
[1] TRUE

R> period(mcE)

[1] 2
```

The example Markov chain found in Mathematica web site (Wolfram Research 2013a) has been used, and is plotted in Figure 4.

```
R> require(matlab)
R> mathematicaMatr <- zeros(5)
R> mathematicaMatr[1,] <- c(0, 1/3, 0, 2/3, 0)
R> mathematicaMatr[2,] <- c(1/2, 0, 0, 0, 1/2)
R> mathematicaMatr[3,] <- c(0, 0, 1/2, 1/2, 0)
R> mathematicaMatr[4,] <- c(0, 0, 1/2, 1/2, 0)
R> mathematicaMatr[5,] <- c(0, 0, 0, 0, 1)
R> statesNames <- letters[1:5]
R> mathematicaMc <- new("markovchain", transitionMatrix = mathematicaMatr,
+                      name = "Mathematica MC", states = statesNames)


Mathematica MC  Markov chain that is composed by:
Closed classes:
c d
```

```
e
Recurrent classes:
{c,d},{e}
Transient classes:
{a,b}
The Markov chain is not irreducible
The absorbing states are: e
```

4.4. First passage time distributions and means

Feres (2007) provides code to compute first passage time (within $1, 2, \ldots, n$ steps) given the initial state to be i. The MATLAB listings translated into R on which the `firstPassage` function is based are:

```
R> .firstpassageKernel <- function(P, i, n){
+    G <- P
+    H <- P[i,]
+    E <- 1 - diag(size(P)[2])
+    for (m in 2:n) {
+      G <- P %*% (G * E)
+      H <- rbind(H, G[i,])
+    }
+    return(H)
+ }
```

We conclude that the probability for the *first* rainy day to be the third one, given that the current state is sunny, is given by:

```
R> firstPassagePdF <- firstPassage(object = mcWeather, state = "sunny",
+                                  n = 10)
R> firstPassagePdF[3, 3]

[1] 0.121
```

To compute the *mean* first passage times, i.e. the expected number of days before it rains given that today is sunny, we can use the `meanFirstPassageTime` function:

```
R> meanFirstPassageTime(mcWeather)

          sunny    cloudy      rain
sunny   0.000000  4.285714  6.666667
cloudy  3.725490  0.000000  5.000000
rain    4.117647  2.857143  0.000000
```

indicating e.g. that the average number of days of sun or cloud before rain is 6.67 if we start counting from a sunny day, and 5 if we start from a cloudy day. Note that we can also specify one or more destination states:

```
R> meanFirstPassageTime(mcWeather,"rain")
```

```
   sunny    cloudy
6.666667 5.000000
```

The implementation follows the matrix solutions by (Grinstead and Snell 2006). We can check the result by averaging the first passage probability density function:

```
R> firstPassagePdF.long <- firstPassage(object = mcWeather, state = "sunny",  n = 100)
R> sum(firstPassagePdF.long[,"rain"] * 1:100)
```

```
[1] 6.666664
```

4.5. Mean recurrence time

The `meanRecurrenceTime` method gives the first mean recurrence time (expected number of steps to go back to a state if it was the initial one) for each recurrent state in the transition probabilities matrix for a DTMC. Let's see an example:

```
R> meanRecurrenceTime(mcWeather)
```

```
   sunny    cloudy      rain
2.156863 3.142857 4.583333
```

Another example, with not all of its states being recurrent:

```
R> recurrentStates(probMc)
```

```
[1] "a" "b" "c" "f" "g" "i"
```

```
R> meanRecurrenceTime(probMc)
```

```
       f        b        g        i        a        c
1.000000 2.555556 2.875000 3.833333 1.500000 3.000000
```

4.6. Absorption probabilities and mean absorption time

We are going to use the Drunkard's random walk from (Grinstead and Snell 2006). We have a drunk person walking through the street. Each move the person does, if they have not arrived to either home (corner 1) or to the bar (corner 5) could be to the left corner or to the right one, with equal probability. In case of arrival to the bar or to home, the person stays there.

```
R> drunkProbs <- matlab::zeros(5, 5)
R> drunkProbs[1,1] <- drunkProbs[5,5] <- 1
R> drunkProbs[2,1] <- drunkProbs[2,3] <- 1/2
R> drunkProbs[3,2] <- drunkProbs[3,4] <- 1/2
R> drunkProbs[4,3] <- drunkProbs[4,5] <- 1/2
R>
R> drunkMc <- new("markovchain", transitionMatrix = drunkProbs)
R> drunkMc


Unnamed Markov chain
 A  5 - dimensional discrete Markov Chain defined by the following states:
 1, 2, 3, 4, 5
 The transition matrix  (by rows)  is defined as follows:
    1   2   3   4   5
1 1.0 0.0 0.0 0.0 0.0
2 0.5 0.0 0.5 0.0 0.0
3 0.0 0.5 0.0 0.5 0.0
4 0.0 0.0 0.5 0.0 0.5
5 0.0 0.0 0.0 0.0 1.0
```

Recurrent (in fact absorbing states) are:

```
R> recurrentStates(drunkMc)


[1] "1" "5"
```

Transient states are the rest:

```
R> transientStates(drunkMc)


[1] "2" "3" "4"
```

The probability of either being absorbed by the bar or by the sofa at home are:

```
R> absorptionProbabilities(drunkMc)


     1    5
2 0.75 0.25
3 0.50 0.50
4 0.25 0.75
```

which means that the probability of arriving home / bar is inversely proportional to the distance to each one.

But we also would like to know how much time does the person take to arrive there, which can be done with `meanAbsorptionTime`:

```
R> meanAbsorptionTime(drunkMc)
```

```
2 3 4
3 4 3
```

So it would take 3 steps to arrive to the destiny if the person is either in the second or fourth corner, and 4 steps in case of being at the same distance from home than to the bar.

4.7. Committor probability

The committor probability tells us the probability to reach a given state before another given. Suppose that we start in a cloudy day, the probabilities of experiencing a rainy day before a sunny one is 0.5:

```
R> committorAB(mcWeather,3,1)
```

```
 sunny cloudy   rain
   0.0    0.5    1.0
```

4.8. Hitting probabilities

Rewriting the system (9) as:

$$
A = \begin{pmatrix}
\begin{array}{c|c|c|c}
A_1 & 0 & \ldots & 0 \\
\hline
0 & A_2 & \ldots & 0 \\
\hline
\vdots & \vdots & \ddots & 0 \\
\hline
0 & 0 & \ldots & A_n
\end{array}
\end{pmatrix}
$$

$$
A_1 = \begin{pmatrix}
-1 & p_{1,2} & p_{1,3} & \cdots & p_{1,n} \\
0 & (p_{2,2}-1) & p_{2,3} & \cdots & p_{2,n} \\
\vdots & \vdots & \vdots & \ddots & \vdots \\
0 & p_{n,2} & p_{n,3} & \cdots & (p_{n,n}-1)
\end{pmatrix}
$$

$$
A_2 = \begin{pmatrix}
(p_{1,1}-1) & 0 & p_{1,3} & \cdots & p_{1,n} \\
p_{2,1} & -1 & p_{2,3} & \cdots & p_{2,n} \\
\vdots & \vdots & \vdots & \ddots & \vdots \\
p_{n,1} & 0 & p_{n,3} & \cdots & (p_{n,n}-1)
\end{pmatrix}
$$

$$
\vdots \qquad\qquad\qquad \vdots
$$

$$
A_n = \begin{pmatrix}
(p_{1,1}-1) & p_{1,2} & p_{1,3} & \cdots & 0 \\
p_{2,1} & (p_{2,2}-1) & p_{2,3} & \cdots & 0 \\
\vdots & \vdots & \vdots & \ddots & \vdots \\
p_{n,1} & p_{n,2} & p_{n,3} & \cdots & -1
\end{pmatrix}
$$

$$X_j = \begin{pmatrix} h_{1,j} \\ h_{2,j} \\ \vdots \\ h_{n,j} \end{pmatrix} \quad C_j = - \begin{pmatrix} p_{1,j} \\ p_{2,j} \\ \vdots \\ p_{n,j} \end{pmatrix}$$

we end up having to solve the block systems:

$$A_j \cdot X_j = C_j \tag{14}$$

Let us imagine the i -th state has transition probabilities: $(0, \ldots, 0, \underset{i)}{1}, 0, \ldots, 0)$. Then that same row would turn into $(0, 0, \ldots, 0)$ for some block, thus obtaining a singular matrix. Another case which may give us problems could be: state i has the following transition probabilities: $(0, \ldots, 0, \underset{j)}{1}, 0, \ldots, 0)$ and the state j has the following transition probabilities: $(0, \ldots, 0, \underset{i)}{1}, 0, \ldots, 0)$. Then when building some blocks we will end up with rows:

$$(0, \ldots, 0, \underset{i)}{-1}, 0, \ldots, 0, \underset{j)}{1}, 0, \ldots, 0)$$
$$(0, \ldots, 0, \underset{i)}{1}, 0, \ldots, 0, \underset{j)}{-1}, 0, \ldots, 0)$$

which are linearly dependent. Our hypothesis is that if we treat the closed communicating classes differently, we *might* delete the linearity in the system. If we have a closed communicating class C_u, then $h_{i,j} = 1$ for all $i, j \in C_u$ and $h_{k,j} = 0$ for all $k \notin C_u$. Then we can set X_u appropriately and solve the other X_v using those values.

The method in charge of that in `markovchain` package is `hittingProbabilities`, which receives a Markov chain and computes the matrix $(h_{ij})_{i,j=1,\ldots,n}$ where $S = \{s_1, \ldots, s_n\}$ is the set of all states of the chain.

For the following chain:

```
R> M <- matlab::zeros(5, 5)
R> M[1,1] <- M[5,5] <- 1
R> M[2,1] <- M[2,3] <- 1/2
R> M[3,2] <- M[3,4] <- 1/2
R> M[4,2] <- M[4,5] <- 1/2
R>
R> hittingTest <- new("markovchain", transitionMatrix = M)
R> hittingProbabilities(hittingTest)

    1     2     3         4   5
1 1.0 0.000 0.000 0.0000000 0.0
2 0.8 0.375 0.500 0.3333333 0.2
3 0.6 0.750 0.375 0.6666667 0.4
4 0.4 0.500 0.250 0.1666667 0.6
5 0.0 0.000 0.000 0.0000000 1.0
```

we want to compute the hitting probabilities. That can be done with:

```
R> hittingProbabilities(hittingTest)
```

```
    1     2     3          4   5
1 1.0 0.000 0.000 0.0000000 0.0
2 0.8 0.375 0.500 0.3333333 0.2
3 0.6 0.750 0.375 0.6666667 0.4
4 0.4 0.500 0.250 0.1666667 0.6
5 0.0 0.000 0.000 0.0000000 1.0
```

In the case of the `mcWeather` Markov chain we would obtain a matrix with all its elements set to 1. That makes sense (and is desirable) since if today is sunny, we expect it would be sunny again at certain point in the time, and the same with rainy weather (that way we assure good harvests):

```
R> hittingProbabilities(mcWeather)
```

```
       sunny cloudy rain
sunny      1      1    1
cloudy     1      1    1
rain       1      1    1
```

5. Statistical analysis

Table 5 lists the functions and methods implemented within the package which help to fit, simulate and predict DTMC.

Function	Purpose
markovchainFit	Function to return fitted Markov chain for a given sequence.
predict	Method to calculate predictions from `markovchain` or `markovchainList` objects.
rmarkovchain	Function to sample from `markovchain` or `markovchainList` objects.

Table 5: The **markovchain** statistical functions.

5.1. Simulation

Simulating a random sequence from an underlying DTMC is quite easy thanks to the function `rmarkovchain`. The following code generates a year of weather states according to `mcWeather` underlying stochastic process.

```
R> weathersOfDays <- rmarkovchain(n = 365, object = mcWeather, t0 = "sunny")
R> weathersOfDays[1:30]
```

```
 [1] "cloudy" "rain"   "rain"   "rain"   "cloudy" "rain"   "rain"   "cloudy"
 [9] "sunny"  "rain"   "cloudy" "cloudy" "cloudy" "rain"   "rain"   "rain"
[17] "cloudy" "rain"   "rain"   "rain"   "rain"   "cloudy" "cloudy" "sunny"
[25] "sunny"  "sunny"  "sunny"  "sunny"  "cloudy" "sunny"
```

Similarly, it is possible to simulate one or more sequences from a non-homogeneous Markov chain, as the following code (applied on CCHC example) exemplifies.

```
R> patientStates <- rmarkovchain(n = 5, object = mcCCRC, t0 = "H",
+                                include.t0 = TRUE)
R> patientStates[1:10,]
```

```
   iteration values
1          1      H
2          1      H
3          1      D
4          1      D
5          1      D
6          2      H
7          2      H
8          2      H
9          2      D
10         2      D
```

Two advance parameters are available to the `rmarkovchain` method which helps you decide which implementation to use. There are four options available : R, R in parallel, C++ and C++ in parallel. Two boolean parameters `useRcpp` and `parallel` will decide which implementation will be used. Default is `useRcpp = TRUE` and `parallel = FALSE` i.e. C++ implementation. The C++ implementation is generally faster than the R implementation. If you have multicore processors then you can take advantage of `parallel` parameter by setting it to TRUE. When both `Rcpp=TRUE` and `parallel=TRUE` the parallelization has been carried out using **RcppParallel** package (Allaire, Francois, Ushey, Vandenbrouck, Geelnard, and Intel 2016).

5.2. Estimation

A time homogeneous Markov chain can be fit from given data. Four methods have been implemented within current version of **markovchain** package: maximum likelihood, maximum likelihood with Laplace smoothing, Bootstrap approach, maximum a posteriori.

Equation 15 shows the maximum likelihood estimator (MLE) of the p_{ij} entry, where the n_{ij} element consists in the number sequences $(X_t = s_i, X_{t+1} = s_j)$ found in the sample, that is

$$\hat{p}_{ij}^{MLE} = \frac{n_{ij}}{\sum\limits_{u=1}^{k} n_{iu}}. \tag{15}$$

Equation (16) shows the `standardError` of the MLE (Skuriat-Olechnowska 2005).

$$SE_{ij} = \frac{\hat{p}_{ij}^{MLE}}{\sqrt{n_{ij}}} \tag{16}$$

```
R> weatherFittedMLE <- markovchainFit(data = weathersOfDays, method = "mle",name = "Weathe
R> weatherFittedMLE$estimate

Weather MLE
 A  3 - dimensional discrete Markov Chain defined by the following states:
 cloudy, rain, sunny
 The transition matrix  (by rows)  is defined as follows:
          cloudy        rain      sunny
cloudy 0.4033613 0.35294118 0.2436975
rain    0.4242424 0.43434343 0.1414141
sunny   0.1917808 0.09589041 0.7123288


R> weatherFittedMLE$standardError


           cloudy        rain      sunny
cloudy 0.05822020 0.05446001 0.04525349
rain    0.06546203 0.06623675 0.03779452
sunny   0.03624317 0.02562779 0.06984958
```

The Laplace smoothing approach is a variation of the MLE, where the n_{ij} is substituted by $n_{ij} + \alpha$ (see Equation 17), being α an arbitrary positive stabilizing parameter.

$$\hat{p}_{ij}^{LS} = \frac{n_{ij} + \alpha}{\sum\limits_{u=1}^{k} (n_{iu} + \alpha)} \tag{17}$$

```
R> weatherFittedLAPLACE <- markovchainFit(data = weathersOfDays,
+                                   method = "laplace", laplacian = 0.01,
+                                   name = "Weather LAPLACE")
R> weatherFittedLAPLACE$estimate

Weather LAPLACE
 A  3 - dimensional discrete Markov Chain defined by the following states:
 cloudy, rain, sunny
 The transition matrix  (by rows)  is defined as follows:
          cloudy        rain      sunny
cloudy 0.4033437 0.35293623 0.2437201
rain    0.4242149 0.43431283 0.1414723
sunny   0.1918099 0.09593919 0.7122509
```

(NOTE: The Confidence Interval option is enabled by default. Remove this option to fasten computations.) Both MLE and Laplace approach are based on the `createSequenceMatrix` functions that returns the raw counts transition matrix.

```
R> createSequenceMatrix(stringchar = weathersOfDays)
```

```
       cloudy rain sunny
cloudy     48   42    29
rain       42   43    14
sunny      28   14   104
```

`stringchar` could contain `NA` values, and the transitions containing `NA` would be ignored.

An issue occurs when the sample contains only one realization of a state (say X_β) which is located at the end of the data sequence, since it yields to a row of zero (no sample to estimate the conditional distribution of the transition). In this case the estimated transition matrix is corrected assuming $p_{\beta,j} = 1/k$, being k the possible states.

Create sequence matrix can also be used to obtain raw count transition matrices from a given $n * 2$ matrix as the following example shows:

```
R> myMatr<-matrix(c("a","b","b","a","a","b","b","b","b","a","a","a","b","a"),ncol=2)
R> createSequenceMatrix(stringchar = myMatr,toRowProbs = TRUE)
```

```
          a         b
a 0.6666667 0.3333333
b 0.5000000 0.5000000
```

A bootstrap estimation approach has been developed within the package in order to provide an indication of the variability of $\hat{p}_{ij}$ estimates. The bootstrap approach implemented within the **markovchain** package follows these steps:

1. bootstrap the data sequences following the conditional distributions of states estimated from the original one. The default bootstrap samples is 10, as specified in `nboot` parameter of `markovchainFit` function.
2. apply MLE estimation on bootstrapped data sequences that are saved in `bootStrapSamples` slot of the returned list.
3. the $p^{BOOTSTRAP}{}_{ij}$ is the average of all $p^{MLE}{}_{ij}$ across the `bootStrapSamples` list, normalized by row. A `standardError` of $p^{\hat{MLE}}{}_{ij}$ estimate is provided as well.

```
R> weatherFittedBOOT <- markovchainFit(data = weathersOfDays,
+                                      method = "bootstrap", nboot = 20)
R> weatherFittedBOOT$estimate
```

```
BootStrap Estimate
 A  3 - dimensional discrete Markov Chain defined by the following states:
 cloudy, rain, sunny
 The transition matrix  (by rows)  is defined as follows:
          cloudy      rain     sunny
cloudy 0.3990505 0.3539818 0.2469677
rain   0.4237268 0.4321171 0.1441561
sunny  0.2034519 0.0922343 0.7043138
```

```
R> weatherFittedBOOT$standardError
```

```
            cloudy        rain       sunny
cloudy 0.008481846 0.005603080 0.007281995
rain   0.008398085 0.008606468 0.005544826
sunny  0.005844212 0.005584338 0.008358063
```

The bootstrapping process can be done in parallel thanks to **RcppParallel** package (Allaire *et al.* 2016). Parallelized implementation is definitely suggested when the data sample size or the required number of bootstrap runs is high.

```
R> weatherFittedBOOTParallel <- markovchainFit(data = weathersOfDays,
+                                  method = "bootstrap", nboot = 200,
+                                  parallel = TRUE)
R> weatherFittedBOOTParallel$estimate
R> weatherFittedBOOTParallel$standardError
```

The parallel bootstrapping uses all the available cores on a machine by default. However, it is also possible to tune the number of threads used. Note that this should be done in R before calling the `markovchainFit` function. For example, the following code will set the number of threads to 4.

```
R> RcppParallel::setNumThreads(2)
```

For more details, please refer to **RcppParallel** web site.

For all the fitting methods, the `logLikelihood` (Skuriat-Olechnowska 2005) denoted in Equation 18 is provided.

$$LLH = \sum_{i,j} n_{ij} * log(p_{ij}) \tag{18}$$

where n_{ij} is the entry of the frequency matrix and p_{ij} is the entry of the transition probability matrix.

```
R> weatherFittedMLE$logLikelihood
```

```
[1] -341.8621
```

```
R> weatherFittedBOOT$logLikelihood
```

```
[1] -341.9377
```

Confidence matrices of estimated parameters (parametric for MLE, non - parametric for Boot-Strap) are available as well. The `confidenceInterval` is provided with the two matrices: `lowerEndpointMatrix` and `upperEndpointMatrix`. The confidence level (CL) is 0.95 by default and can be given as an argument of the function `markovchainFit`. This is used to obtain

the standard score (z-score). From classical inference theory, if ci is the level of confidence required assuming normal distribution the $zscore(ci)$ solves $\Phi\left(1 - \left(\frac{1-ci}{2}\right)\right)$ Equations 19 and 20 (Skuriat-Olechnowska 2005) show the `confidenceInterval` of a fitting. Note that each entry of the matrices is bounded between 0 and 1.

$$LowerEndpoint_{ij} = p_{ij} - zscore(CL) * SE_{ij} \qquad (19)$$
$$UpperEndpoint_{ij} = p_{ij} + zscore(CL) * SE_{ij} \qquad (20)$$

```
R> weatherFittedMLE$confidenceInterval


NULL


R> weatherFittedBOOT$confidenceInterval


$confidenceLevel
[1] 0.95


$lowerEndpointMatrix
          cloudy       rain      sunny
cloudy 0.3850991 0.34476555 0.2349899
rain   0.4099132 0.41796070 0.1350357
sunny  0.1938390 0.08304888 0.6905660


$upperEndpointMatrix
          cloudy      rain      sunny
cloudy 0.4130019 0.3631980 0.2589455
rain   0.4375404 0.4462735 0.1532766
sunny  0.2130648 0.1014197 0.7180616
```

A special function, `multinomialConfidenceIntervals`, has been written in order to obtain multinomial wise confidence intervals. The code has been based on and Rcpp translation of package's **MultinomialCI** functions Villacorta (2012) that were themselves based on the Sison and Glaz (1995) paper.

```
R> multinomialConfidenceIntervals(transitionMatrix =
+        weatherFittedMLE$estimate@transitionMatrix,
+        countsTransitionMatrix = createSequenceMatrix(weathersOfDays))


$confidenceLevel
[1] 0.95


$lowerEndpointMatrix
          cloudy       rain      sunny
cloudy 0.3109244 0.26050420 0.15126050
rain   0.3232323 0.33333333 0.04040404
```

```
sunny   0.1232877 0.02739726 0.64383562


$upperEndpointMatrix
          cloudy      rain     sunny
cloudy 0.5031240 0.4527039 0.3434602
rain   0.5284189 0.5385200 0.2455907
sunny  0.2655683 0.1696779 0.7861162
```

The functions for fitting DTMC have mostly been rewritten in C++ using **Rcpp** Eddelbuettel (2013) since version 0.2.

It is also possible to fit a DTMC object from `matrix` or `data.frame` objects as shown in following code.

```
R> data(holson)
R> singleMc<-markovchainFit(data=holson[,2:12],name="holson")
```

The same applies for `markovchainList` (output length has been limited).

```
R> mcListFit<-markovchainListFit(data=holson[,2:6],name="holson")
R> mcListFit$estimate
```

```
holson  list of Markov chain(s)
Markovchain  1
Unnamed Markov chain
 A  1 - dimensional discrete Markov Chain defined by the following states:
 1
 The transition matrix  (by rows)  is defined as follows:
  1
1 1


Markovchain  2
Unnamed Markov chain
 A  1 - dimensional discrete Markov Chain defined by the following states:
 1
 The transition matrix  (by rows)  is defined as follows:
  1
1 1


Markovchain  3
Unnamed Markov chain
 A  2 - dimensional discrete Markov Chain defined by the following states:
...
```

Finally, given a `list` object, it is possible to fit a `markovchain` object or to obtain the raw transition matrix.

```
R> c1<-c("a","b","a","a","c","c","a")
R> c2<-c("b")
R> c3<-c("c","a","a","c")
R> c4<-c("b","a","b","a","a","c","b")
R> c5<-c("a","a","c",NA)
R> c6<-c("b","c","b","c","a")
R> mylist<-list(c1,c2,c3,c4,c5,c6)
R> mylistMc<-markovchainFit(data=mylist)
R> mylistMc

$estimate
MLE Fit
 A  3 - dimensional discrete Markov Chain defined by the following states:
 a, b, c
 The transition matrix  (by rows)  is defined as follows:
    a         b         c
a 0.4 0.2000000 0.4000000
b 0.6 0.0000000 0.4000000
c 0.5 0.3333333 0.1666667


$standardError
          a         b         c
a 0.2000000 0.1414214 0.2000000
b 0.3464102 0.0000000 0.2828427
c 0.2886751 0.2357023 0.1666667

$confidenceLevel
[1] 0.95

$lowerEndpointMatrix
            a b           c
a 0.008007122 0 0.008007122
b 0.000000000 0 0.000000000
c 0.000000000 0 0.000000000

$upperEndpointMatrix
          a         b         c
a 0.7919929 0.4771808 0.7919929
b 1.0000000 0.0000000 0.9543616
c 1.0000000 0.7953014 0.4933274
```

The same works for `markovchainFitList`.

```
R> markovchainListFit(data=mylist)

$estimate
  list of Markov chain(s)
```

```
Markovchain  1
Unnamed Markov chain
 A  3 - dimensional discrete Markov Chain defined by the following states:
 a, b, c
 The transition matrix  (by rows)  is defined as follows:
    a   b   c
a 0.5 0.5 0.0
b 0.5 0.0 0.5
c 1.0 0.0 0.0

Markovchain  2
Unnamed Markov chain
 A  3 - dimensional discrete Markov Chain defined by the following states:
...
```

If any transition contains `NA`, it will be ignored in the results as the above example showed.

5.3. Prediction

The n-step forward predictions can be obtained using the `predict` methods explicitly written for `markovchain` and `markovchainList` objects. The prediction is the mode of the conditional distribution of X_{t+1} given $X_t = s_j$, being s_j the last realization of the DTMC (homogeneous or non-homogeneous).

Predicting from a markovchain object

The 3-days forward predictions from `markovchain` object can be generated as follows, assuming that the last two days were respectively "cloudy" and "sunny".

```
R> predict(object = weatherFittedMLE$estimate, newdata = c("cloudy", "sunny"),
+        n.ahead = 3)

[1] "sunny" "sunny" "sunny"
```

Predicting from a markovchainList object

Given an initial two years health status, the 5-year ahead prediction of any CCRC guest is

```
R> predict(mcCCRC, newdata = c("H", "H"), n.ahead = 5)

[1] "H" "D" "D"
```

The prediction has stopped at time sequence since the underlying non-homogeneous Markov chain has a length of four. In order to continue five years ahead, the `continue=TRUE` parameter setting makes the `predict` method keeping to use the last `markovchain` in the sequence list.

```
R> predict(mcCCRC, newdata = c("H", "H"), n.ahead = 5, continue = TRUE)
```

```
[1] "H" "D" "D" "D" "D"
```

5.4. Statistical Tests

In this section, we describe the statistical tests: assessing the Markov property (`verifyMarkovProperty`), the order (`assessOrder`), the stationary (`assessStationarity`) of a Markov chain sequence, and the divergence test for empirically estimated transition matrices (`divergenceTest`). Most of such tests are based on the χ^2 statistics. Relevant references are Kullback *et al.* (1962) and Anderson and Goodman (1957).

All such tests have been designed for small samples, since it is easy to detect departures from Markov property as long as the sample size increases. In addition, the accuracy of the statistical inference functions has been questioned and will be thoroughly investigated in future versions of the package.

Assessing the Markov property of a Markov chain sequence

The `verifyMarkovProperty` function verifies whether the Markov property holds for the given chain. The test implemented in the package looks at triplets of successive observations. If $x_1, x_2, \ldots, x_N$ is a set of observations and n_{ijk} is the number of times t ($1 \leq t \leq N-2$) such that $x_t = i, x_{t+1} = j, x_{x+2} = k$, then if the Markov property holds n_{ijk} follows a Binomial distribution with parameters n_{ij} and p_{jk}. A classical χ^2 test can check this distributional assumption, since $\sum_i \sum_j \sum_k \frac{(n_{ijk} - n_{ij}\hat{p_{jk}})^2}{n_{ij}\hat{p_{jk}}} \sim \chi^2(q)$ where q is the number of degrees of freedom. The number of degrees of freedom q of the distribution of χ^2 is given by the formula r-q+s-1, where: s denotes the number of states i in the state space such that n_{i} > 0 q denotes the number of pairs (i, j) for which n_{ij} > 0 and r denotes the number of triplets (i, j, k) for which n_{ij}n_{jk} > 0

```
R> sample_sequence<-c("a", "b", "a", "a", "a", "a", "b", "a", "b", "a",
+                     "b", "a", "a", "b", "b", "b", "a")
R> verifyMarkovProperty(sample_sequence)


Testing markovianity property on given data sequence
Chi - square statistic is: 0.28
Degrees of freedom are: 5
And corresponding p-value is: 0.9980024
```

Assessing the order of a Markov chain sequence

The `assessOrder` function checks whether the given chain is of first order or of second order. For each possible present state, we construct a contingency table of the frequency of the future state for each past to present state transition as shown in Table 6.

Using the table, the function performs the χ^2 test by calling the `chisq.test` function. This test returns a list of the chi-squared value and the p-value. If the p-value is greater than the given significance level, we cannot reject the hypothesis that the sequence is of first order.

past	present	future a	future b
a	a	2	2
b	a	2	2

Table 6: Contingency table to assess the order for the present state a.

```
R> data(rain)
R> assessOrder(rain$rain)


Warning in assessOrder(rain$rain): The accuracy of the statistical inference
functions has been questioned. It will be thoroughly investigated in future
versions of the package.


The assessOrder test statistic is:  26.09575
The Chi-Square d.f. are:  12
The p-value is:  0.01040395
```

Assessing the stationarity of a Markov chain sequence

The `assessStationarity` function assesses if the transition probabilities of the given chain change over time. To be more specific, the chain is stationary if the following condition meets.

$$p_{ij}(t) = p_{ij} \quad \text{for all} \ \ t \tag{21}$$

For each possible state, we construct a contingency table of the estimated transition probabilities over time as shown in Table 7.

time (t)	probability of transition to a	probability of transition to b
1	0	1
2	0	1
.	.	.
.	.	.
.	.	.
16	0.44	0.56

Table 7: Contingency table to assess the stationarity of the state a.

Using the table, the function performs the χ^2 test by calling the `chisq.test` function. This test returns a list of the chi-squared value and the p-value. If the p-value is greater than the given significance level, we cannot reject the hypothesis that the sequence is stationary.

```
R> assessStationarity(rain$rain, 10)


Warning in assessStationarity(rain$rain, 10): The accuracy of the statistical
inference functions has been questioned. It will be thoroughly investigated in
future versions of the package.
```

```
Warning in chisq.test(mat): Chi-squared approximation may be incorrect

Warning in chisq.test(mat): Chi-squared approximation may be incorrect

Warning in chisq.test(mat): Chi-squared approximation may be incorrect

The assessStationarity test statistic is:  4.181815
The Chi-Square d.f. are:  54
The p-value is:  1
```

Divergence tests for empirically estimated transition matrices

This section discusses tests developed to verify whether:

1. An empirical transition matrix is consistent with a theoretical one.

2. Two or more empirical transition matrices belongs to the same DTMC.

The first test is implemented by the `verifyEmpiricalToTheoretical` function. Being f_{ij} the raw transition count, Kullback *et al.* (1962) shows that $2 * \sum_{i=1}^{r} \sum_{j=1}^{r} f_{ij} \ln \frac{f_{ij}}{f_{i.} P(E_j|E_i)} \sim \chi^2 (r * (r - 1))$. The following example is taken from Kullback *et al.* (1962):

```
R> sequence<-c(0,1,2,2,1,0,0,0,0,0,0,1,2,2,2,1,0,0,1,0,0,0,0,0,0,1,1,
+ 2,0,0,2,1,1,0,0,0,0,0,0,0,0,0,0,0,0,0,1,1,1,1,0,0,0,0,2,1,0,
+ 0,2,1,0,0,0,0,0,0,1,1,1,2,2,0,0,2,1,1,1,1,2,1,1,1,1,1,1,1,1,0,2,
+ 0,1,1,0,0,0,1,2,2,0,0,0,0,0,0,2,2,2,1,1,1,1,0,1,1,1,1,0,0,2,1,1,
+ 0,0,0,0,0,2,2,1,1,1,1,1,2,1,2,0,0,0,1,2,2,2,0,0,0,1,1)
R> mc=matrix(c(5/8,1/4,1/8,1/4,1/2,1/4,1/4,3/8,3/8),byrow=TRUE, nrow=3)
R> rownames(mc)<-colnames(mc)<-0:2; theoreticalMc<-as(mc, "markovchain")
R> verifyEmpiricalToTheoretical(data=sequence,object=theoreticalMc)

Testing whether the
   0  1  2
0 51 11  8
1 12 31  9
2  6 11 10
transition matrix is compatible with
      0     1     2
0 0.625 0.250 0.125
1 0.250 0.500 0.250
2 0.250 0.375 0.375
[1] "theoretical transition matrix"
ChiSq statistic is 6.551795 d.o.f are 6 corresponding p-value is 0.3642899

$statistic
      0
```

```
6.551795
```

```
$dof
[1] 6
```

```
$pvalue
         0
0.3642899
```

The second one is implemented by the `verifyHomogeneity` function, inspired by (Kullback *et al.* 1962, section 9). Assuming that $i = 1, 2, \ldots, s$ DTMC samples are available and that the cardinality of the state space is r it verifies whether the s chains belongs to the same unknown one. Kullback *et al.* (1962) shows that its test statistics follows a chi-square law, $2 * \sum_{i=1}^{s} \sum_{j=1}^{r} \sum_{k=1}^{r} f_{ijk} \ln \frac{n * f_{ijk}}{f_{i..} f_{.jk}} \sim \chi^2 \left(r * (r - 1) \right)$. Also the following example is taken from Kullback *et al.* (1962):

```
R> data(kullback)
R> verifyHomogeneity(inputList=kullback,verbose=TRUE)
```

```
Warning in verifyHomogeneity(inputList = kullback, verbose = TRUE): The
accuracy of the statistical inference functions has been questioned. It will be
thoroughly investigated in future versions of the package.
```

```
Testing homogeneity of DTMC underlying input list
ChiSq statistic is 275.9963 d.o.f are 35 corresponding p-value is 0
```

```
$statistic
[1] 275.9963
```

```
$dof
[1] 35
```

```
$pvalue
[1] 0
```

5.5. Continuous Times Markov Chains

Intro

The **markovchain** package provides functionality for continuous time Markov chains (CTMCs). CTMCs are a generalization of discrete time Markov chains (DTMCs) in that we allow time to be continuous. We assume a finite state space S (for an infinite state space wouldn't fit in memory). We can think of CTMCs as Markov chains in which state transitions can happen at any time.

More formally, we would like our CTMCs to satisfy the following two properties:

- The Markov property - let $F_{X(s)}$ denote the information about X up to time s. Let $j \in S$ and $s \leq t$. Then, $P(X(t) = j | F_{X(s)}) = P(X(t) = j | X(s))$.
- Time homogeneity - $P(X(t) = j | X(s) = k) = P(X(t - s) = j | X(0) = k)$.

If both the above properties are satisfied, it is referred to as a time-homogeneous CTMC. If a transition occurs at time t, then $X(t)$ denotes the new state and $X(t) \neq X(t-)$.

Now, let $X(0) = x$ and let T_x be the time a transition occurs from this state. We are interested in the distribution of T_x. For $s, t \geq 0$, it can be shown that $P(T\_x > s+t \mid T\_x > s) = P(T\_x > t)$

This is the memory less property that only the exponential random variable exhibits. Therefore, this is the sought distribution, and each state $s \in S$ has an exponential holding parameter $\lambda(s)$. Since $\mathrm{E}T_x = \frac{1}{\lambda(x)}$, higher the rate $\lambda(x)$, smaller the expected time of transitioning out of the state x.

However, specifying this parameter alone for each state would only paint an incomplete picture of our CTMC. To see why, consider a state x that may transition to either state y or z. The holding parameter enables us to predict when a transition may occur if we start off in state x, but tells us nothing about which state will be next.

To this end, we also need transition probabilities associated with the process, defined as follows (for $y \neq x$) - $p_{xy} = P(X(T_s) = y | X(0) = x)$. Note that $\sum_{y \neq x} p_{xy} = 1$. Let Q denote this transition matrix ($Q_{ij} = p_{ij}$). What is key here is that T_x and the state y are independent random variables. Let's define $\lambda(x, y) = \lambda(x)p_{xy}$

We now look at Kolmogorov's backward equation. Let's define $P_{ij}(t) = P(X(t) = j | X(0) = i)$ for $i, j \in S$. The backward equation is given by (it can be proved) $P_{ij}(t) = \delta_{ij}e^{-\lambda(i)t} + \int_0^t \lambda(i)e^{-\lambda(i)t} \sum_{k \neq i} Q_{ik}P_{kj}(t - s)ds$. Basically, the first term is non-zero if and only if $i = j$ and represents the probability that the first transition from state i occurs after time t. This would mean that at t, the state is still i. The second term accounts for any transitions that may occur before time t and denotes the probability that at time t, when the smoke clears, we are in state j.

This equation can be represented compactly as follows $P'(t) = AP(t)$ where A is the *generator* matrix.

$$A(i, j) = \begin{cases} \lambda(i, j) & \text{if } i \neq j \\ -\lambda(i) & \text{else.} \end{cases}$$

Observe that the sum of each row is 0. A CTMC can be completely specified by the generator matrix.

Stationary Distributions

The following theorem guarantees the existence of a unique stationary distribution for CTMCs. Note that $X(t)$ being irreducible and recurrent is the same as $X_n(t)$ being irreducible and recurrent.

Suppose that $X(t)$ is irreducible and recurrent. Then $X(t)$ has an invariant measure η, which is unique up to multiplicative factors. Moreover, for each $k \in S$, we have

$$\eta_k = \frac{\pi_k}{\lambda(k)}$$

where π is the unique invariant measure of the embedded discrete time Markov chain Xn. Finally, η satisfies

$$0 < \eta_j < \infty, \forall j \in S$$

and if $\sum_i \eta_i < \infty$ then η can be normalized to get a stationary distribution.

Estimation

Let the data set be $D = \{(s_0, t_0), (s_1, t_1), ..., (s_{N-1}, t_{N-1})\}$ where $N = |D|$. Each s_i is a state from the state space S and during the time $[t_i, t_{i+1}]$ the chain is in state s_i. Let the parameters be represented by $\theta = \{\lambda, P\}$ where λ is the vector of holding parameters for each state and P the transition matrix of the embedded discrete time Markov chain.

Then the probability is given by

$$Pr(D|\theta) \propto \lambda(s_0)e^{-\lambda(s_0)(t_1 - t_0)}Pr(s_1|s_0) \cdot ... \cdot \lambda(s_{N-2})e^{-\lambda(s_{N-2})(t_{N-1} - t_{N-2})}Pr(s_{N-1}|s_{N-2})$$

Let $n(j|i)$ denote the number of i->j transitions in D, and $n(i)$ the number of times s_i occurs in D. Let $t(s_i)$ denote the total time the chain spends in state s_i.

Then the MLEs are given by

$$\hat{\lambda(s)} = \frac{n(s)}{t(s)}, \hat{Pr}(j|i) = \frac{n(j|i)}{n(i)}$$

Expected Hitting Time

The package provides a function `ExpectedTime` to calculate average hitting time from one state to another. Let the final state be j, then for every state $i \in S$, where S is the set of all states and holding time $q_i > 0$ for every $i \neq j$. Assuming the conditions to be true, expected hitting time is equal to minimal non-negative solution vector p to the system of linear equations:

$$\begin{cases} p_k = 0 & k = j \\ -\sum_{l \in I} q_{kl}p_k = 1 & k \neq j \end{cases} \tag{22}$$

Probability at time t

The package provides a function `probabilityatT` to calculate probability of every state according to given `ctmc` object. Here we use Kolmogorov's backward equation $P(t) = P(0)e^{tQ}$ for $t \geq 0$ and $P(0) = I$. Here $P(t)$ is the transition function at time t. The value $P(t)[i][j]$ at time $P(t)$ describes the probability of the state at time t to be equal to j if it was equal to i at time $t = 0$. It takes care of the case when `ctmc` object has a generator represented by columns. If initial state is not provided, the function returns the whole transition matrix $P(t)$.

Examples

To create a CTMC object, you need to provide a valid generator matrix, say Q. The CTMC object has the following slots - states, generator, by row, name (look at the documentation object for further details). Consider the following example in which we aim to model the transition of a molecule from the σ state to the σ^* state. When in the former state, if it absorbs sufficient energy, it can make the jump to the latter state and remains there for some time before transitioning back to the original state. Let us model this by a CTMC:

```
R> energyStates <- c("sigma", "sigma_star")
R> byRow <- TRUE
R> gen <- matrix(data = c(-3, 3,
+                          1, -1), nrow = 2,
+           byrow = byRow, dimnames = list(energyStates, energyStates))
R> molecularCTMC <- new("ctmc", states = energyStates,
+               byrow = byRow, generator = gen,
+               name = "Molecular Transition Model")
```

To generate random CTMC transitions, we provide an initial distribution of the states. This must be in the same order as the dimnames of the generator. The output can be returned either as a list or a data frame.

```
R> statesDist <- c(0.8, 0.2)
R> rctmc(n = 3, ctmc = molecularCTMC, initDist = statesDist, out.type = "df", include.T0 =
```

```
      states              time
1 sigma_star 0.103381332786133
2      sigma    1.863359310425
3 sigma_star  2.45358102090009
```

n represents the number of samples to generate. There is an optional argument T for `rctmc`. It represents the time of termination of the simulation. To use this feature, set n to a very high value, say `Inf` (since we do not know the number of transitions before hand) and set T accordingly.

```
R> statesDist <- c(0.8, 0.2)
R> rctmc(n = Inf, ctmc = molecularCTMC, initDist = statesDist, T = 2)
```

```
[[1]]
[1] "sigma"      "sigma_star" "sigma"      "sigma_star" "sigma"

[[2]]
[1] 0.0000000 0.0426797 0.4407189 1.7328868 1.7932432
```

To obtain the stationary distribution simply invoke the `steadyStates` function

```
R> steadyStates(molecularCTMC)
```

```
      sigma sigma_star
[1,]  0.25       0.75
```

For fitting, use the `ctmcFit` function. It returns the MLE values for the parameters along with the confidence intervals.

```
R> data <- list(c("a", "b", "c", "a", "b", "a", "c", "b", "c"),
+              c(0, 0.8, 2.1, 2.4, 4, 5, 5.9, 8.2, 9))
R> ctmcFit(data)

$estimate
An object of class "ctmc"
Slot "states":
[1] "a" "b" "c"

Slot "byrow":
[1] TRUE

Slot "generator":
           a          b          c
a -0.9090909  0.6060606  0.3030303
b  0.3225806 -0.9677419  0.6451613
c  0.3846154  0.3846154 -0.7692308

Slot "name":
[1] ""


$errors
$errors$dtmcConfidenceInterval
$errors$dtmcConfidenceInterval$confidenceLevel
[1] 0.95

$errors$dtmcConfidenceInterval$lowerEndpointMatrix
  a b c
a 0 0 0
b 0 0 0
c 0 0 0

$errors$dtmcConfidenceInterval$upperEndpointMatrix
          a b         c
a 0.0000000 1 0.9866548
b 0.9866548 0 1.0000000
c 1.0000000 1 0.0000000


$errors$lambdaConfidenceInterval
```

```
$errors$lambdaConfidenceInterval$lowerEndpointVector
[1] 0.04576665 0.04871934 0.00000000

$errors$lambdaConfidenceInterval$upperEndpointVector
[1]  0.04576665  0.04871934 -0.12545166
```

One approach to obtain the generator matrix is to apply the `logm` function from the **expm** package on a transition matrix. Numeric issues arise, see Israel, Rosenthal, and Wei (2001). For example, applying the standard `method` ('Higham08') on `mcWeather` raises an error, whilst the alternative method (eigenvalue decomposition) is OK. The following code estimates the generator matrix of the `mcWeather` transition matrix.

```
R> mcWeatherQ <- expm::logm(mcWeather@transitionMatrix,method='Eigen')
R> mcWeatherQ


            sunny      cloudy        rain
sunny   -0.863221    2.428723   -1.565502
cloudy   4.284592  -20.116312   15.831720
rain    -4.414019   24.175251  -19.761232
```

Therefore, the "half - day" transition probability for mcWeather DTMC is

```
R> mcWeatherHalfDayTM <- expm::expm(mcWeatherQ*.5)
R> mcWeatherHalfDay <- new("markovchain",transitionMatrix=mcWeatherHalfDayTM,name="Half Da
R> mcWeatherHalfDay

Half Day Weather Transition Matrix
 A  3 - dimensional discrete Markov Chain defined by the following states:
 sunny, cloudy, rain
 The transition matrix  (by rows)  is defined as follows:
             sunny     cloudy        rain
sunny   0.81598647 0.1420068 0.04200677
cloudy 0.21970167 0.4401492 0.34014916
rain    0.07063048 0.5146848 0.41468476
```

The **ctmcd** package (Pfeuffer 2017) provides various functions to estimate the generator matrix (GM) of a CTMC process using different methods. The following code provides a way to join **markovchain** and **ctmcd** computations.

```
R> require(ctmcd)

Loading required package: ctmcd


R> require(expm)

Loading required package: expm
```

```
Loading required package: Matrix


Attaching package: 'expm'

The following object is masked from 'package:Matrix':

    expm
```

```
R> #defines a function to transform a GM into a TM
R> gm_to_markovchain<-function(object, t=1) {
+   if(!(class(object) %in% c("gm","matrix","Matrix")))
+     stop("Error! Expecting either a matrix or a gm object")
+   if ( class(object) %in% c("matrix","Matrix")) generator_matrix<-object else generator_
+   #must add importClassesFrom("markovchain",markovchain) in the NAMESPACE
+   #must add importFrom(expm, "expm")
+   transitionMatrix<-expm(generator_matrix*t)
+   out<-as(transitionMatrix,"markovchain")
+   return(out)
+ }
R> #loading ctmcd dataset
R> data(tm_abs)
R> gm0=matrix(1,8,8) #initializing
R> diag(gm0)=0
R> diag(gm0)=-rowSums(gm0)
R> gm0[8,]=0
R> gmem=gm(tm_abs,te=1,method="EM",gmguess=gm0) #estimating GM
R> mc_at_2=gm_to_markovchain(object=gmem, t=2) #converting to TM at time 2
```

5.6. Pseudo - Bayesian Estimation

Hu, Kiesel, and Perraudin (2002) shows an empirical quasi-Bayesian method to estimate transition matrices, given an empirical $\hat{P}$ transition matrix (estimated using the classical approach) and an a - priori estimate Q. In particular, each row of the matrix is estimated using the linear combination $\alpha \cdot Q + (1 - 1alpha) \cdot P$, where α is defined for each row as Equation 23 shows

$$\begin{cases} \hat{\alpha}_i = \frac{\hat{K}_i}{v(i)+\hat{K}_i} \\ \hat{K}_i = \frac{v(i)^2 - \sum_j Y_{ij}^2}{\sum_j (Y_{ij} - v(i)*q_{ij})^2} \end{cases} \tag{23}$$

The following code returns the pseudo Bayesian estimate of the transition matrix:

```
R> pseudoBayesEstimator <- function(raw, apriori){
+   v_i <- rowSums(raw)
+   K_i <- numeric(nrow(raw))
```

```
+    sumSquaredY <- rowSums(raw^2)
+    #get numerator
+    K_i_num <- v_i^2-sumSquaredY
+    #get denominator
+    VQ <- matrix(0,nrow= nrow(apriori),ncol=ncol(apriori))
+    for (i in 1:nrow(VQ)) {
+      VQ[i,]<-v_i[i]*apriori[i,]
+    }
+
+    K_i_den<-rowSums((raw - VQ)^2)
+
+    K_i <- K_i_num/K_i_den
+
+    #get the alpha vector
+    alpha <- K_i / (v_i+K_i)
+
+    #empirical transition matrix
+    Emp<-raw/rowSums(raw)
+
+    #get the estimate
+    out<-matrix(0, nrow= nrow(raw),ncol=ncol(raw))
+    for (i in 1:nrow(out)) {
+      out[i,]<-alpha[i]*apriori[i,]+(1-alpha[i])*Emp[i,]
+    }
+    return(out)
+ }
```

We then apply it to the weather example:

```
R> trueMc<-as(matrix(c(0.1, .9,.7,.3),nrow = 2, byrow = 2),"markovchain")
R> aprioriMc<-as(matrix(c(0.5, .5,.5,.5),nrow = 2, byrow = 2),"markovchain")
R>
R> smallSample<-rmarkovchain(n=20,object = trueMc)
R> smallSampleRawTransitions<-createSequenceMatrix(stringchar = smallSample)
R> pseudoBayesEstimator(
+    raw = smallSampleRawTransitions,
+    apriori = aprioriMc@transitionMatrix
+ ) - trueMc@transitionMatrix

          s1         s2
s1 -0.100000 0.100000
s2 -0.136019 0.136019


R> biggerSample<-rmarkovchain(n=100,object = trueMc)
R> biggerSampleRawTransitions<-createSequenceMatrix(stringchar = biggerSample)
R> pseudoBayesEstimator(
+    raw = biggerSampleRawTransitions,
```

```
+   apriori = aprioriMc@transitionMatrix
+ ) - trueMc@transitionMatrix


            s1            s2
s1  0.02240738 -0.02240738
s2 -0.02107810  0.02107810


R> bigSample<-rmarkovchain(n=1000,object = trueMc)
R> bigSampleRawTransitions<-createSequenceMatrix(stringchar = bigSample)
R> pseudoBayesEstimator(
+   raw = bigSampleRawTransitions,
+   apriori = aprioriMc@transitionMatrix
+ ) - trueMc@transitionMatrix


            s1            s2
s1 0.011686151 -0.011686151
s2 0.002490415 -0.002490415
```

5.7. Bayesian Estimation

The **markovchain** package provides functionality for maximum a posteriori (MAP) estimation of the chain parameters (at the time of writing this document, only first order models are supported) by Bayesian inference. It also computes the probability of observing a new data set, given a (different) data set. This vignette provides the mathematical description for the methods employed by the package.

Notation and set-up

The data is denoted by D, the model parameters (transition matrix) by θ. The object of interest is $P(\theta|D)$ (posterior density). $\mathcal{A}$ represents an alphabet class, each of whose members represent a state of the chain. Therefore

$$D = s_0 s_1 ... s_{N-1}, s_t \in \mathcal{A}$$

where N is the length of the data set. Also,

$$\theta = \{p(s|u), s \in \mathcal{A}, u \in \mathcal{A}\}$$

where $\sum_{s \in \mathcal{A}} p(s|u) = 1$ for each $u \in \mathcal{A}$.

Our objective is to find θ which maximizes the posterior. That is, if our solution is denoted by $\hat{\theta}$, then

$$\hat{\theta} = \underset{\theta}{argmax} P(\theta|D)$$

where the search space is the set of right stochastic matrices of dimension $|\mathcal{A}| x |\mathcal{A}|$.

$n(u, s)$ denotes the number of times the word us occurs in D and $n(u) = \sum_{s \in \mathcal{A}} n(u, s)$. The hyper-parameters are similarly denoted by $\alpha(u, s)$ and $\alpha(u)$ respectively.

Methods

Given D, its likelihood conditioned on the observed initial state in D is given by

$$P(D|\theta) = \prod_{s \in \mathcal{A}} \prod_{u \in \mathcal{A}} p(s|u)^{n(u,s)}$$

Conjugate priors are used to model the prior $P(\theta)$. The reasons are two fold:

1. Exact expressions can be derived for the MAP estimates, expectations and even variances
2. Model order selection/comparison can be implemented easily (available in a future release of the package)

The hyper-parameters determine the form of the prior distribution, which is a product of Dirichlet distributions

$$P(\theta) = \prod_{u \in \mathcal{A}} \left\{ \frac{\Gamma(\alpha(u))}{\prod_{s \in \mathcal{A}} \Gamma(\alpha(u, s))} \prod_{s \in \mathcal{A}} p(s|u)^{\alpha(u,s))-1} \right\}$$

where $\Gamma(.)$ is the Gamma function. The hyper-parameters are specified using the `hyperparam` argument in the `markovchainFit` function. If this argument is not specified, then a default value of 1 is assigned to each hyper-parameter resulting in the prior distribution of each chain parameter to be uniform over $[0, 1]$.

Given the likelihood and the prior as described above, the evidence $P(D)$ is simply given by

$$P(D) = \int P(D|\theta) P(\theta) d\theta$$

which simplifies to

$$P(D) = \prod_{u \in \mathcal{A}} \left\{ \frac{\Gamma(\alpha(u))}{\prod_{s \in \mathcal{A}} \Gamma(\alpha(u, s))} \frac{\prod_{s \in \mathcal{A}} \Gamma(n(u, s) + \alpha(u, s))}{\Gamma(\alpha(u) + n(u))} \right\}$$

Using Bayes' theorem, the posterior now becomes (thanks to the choice of conjugate priors)

$$P(\theta|D) = \prod_{u \in \mathcal{A}} \left\{ \frac{\Gamma(n(u) + \alpha(u))}{\prod_{s \in \mathcal{A}} \Gamma(n(u, s) + \alpha(u, s))} \prod_{s \in \mathcal{A}} p(s|u)^{n(u,s)+\alpha(u,s))-1} \right\}$$

Since this is again a product of Dirichlet distributions, the marginal distribution of a particular parameter $P(s|u)$ of our chain is given by

$$P(s|u) \sim Beta(n(u, s) + \alpha(u, s), n(u) + \alpha(u) - n(u, s) - \alpha(u, s))$$

Thus, the MAP estimate $\hat{\theta}$ is given by

$$\hat{\theta} = \left\{ \frac{n(u, s) + \alpha(u, s) - 1}{n(u) + \alpha(u) - |\mathcal{A}|}, s \in \mathcal{A}, u \in \mathcal{A} \right\}$$

The function also returns the expected value, given by

$$\mathrm{E}_{\mathrm{post}}p(s|u) = \left\{\frac{n(u,s) + \alpha(u,s)}{n(u) + \alpha(u)}, s \in \mathcal{A}, u \in \mathcal{A}\right\}$$

The variance is given by

$$\mathrm{Var}_{\mathrm{post}}p(s|u) = \frac{n(u,s) + \alpha(u,s)}{(n(u) + \alpha(u))^2}\frac{n(u) + \alpha(u) - n(u,s) - \alpha(u,s)}{n(u) + \alpha(u) + 1}$$

The square root of this quantity is the standard error, which is returned by the function. The confidence intervals are constructed by computing the inverse of the beta integral.

Predictive distribution

Given the old data set, the probability of observing new data is $P(D'|D)$ where D' is the new data set. Let $m(u,s), m(u)$ denote the corresponding counts for the new data. Then,

$$P(D'|D) = \int P(D'|\theta)P(\theta|D)d\theta$$

We already know the expressions for both quantities in the integral and it turns out to be similar to evaluating the evidence

$$P(D'|D) = \prod_{u \in \mathcal{A}}\left\{\frac{\Gamma(\alpha(u))}{\prod_{s \in \mathcal{A}}\Gamma(\alpha(u,s))}\frac{\prod_{s \in \mathcal{A}}\Gamma(n(u,s) + m(u,s) + \alpha(u,s))}{\Gamma(\alpha(u) + n(u) + m(u))}\right\}$$

Choosing the hyper-parameters

The hyper parameters model the shape of the parameters' prior distribution. These must be provided by the user. The package offers functionality to translate a given prior belief transition matrix into the hyper-parameter matrix. It is assumed that this belief matrix corresponds to the mean value of the parameters. Since the relation

$$\mathrm{E}_{\mathrm{prior}}p(s|u) = \frac{\alpha(u,s)}{\alpha(u)}$$

holds, the function accepts as input the belief matrix as well as a scaling vector (serves as a proxy for $\alpha(.)$) and proceeds to compute $\alpha(.,.)$.

Alternatively, the function accepts a data sample and infers the hyper-parameters from it. Since the mode of a parameter (with respect to the prior distribution) is proportional to one less than the corresponding hyper-parameter, we set

$$\alpha(u,s) - 1 = m(u,s)$$

where $m(u,s)$ is the $u \rightarrow s$ transition count in the data sample. This is regarded as a 'fake count' which helps $\alpha(u,s)$ to reflect knowledge of the data sample.

Usage and examples

```
R> weatherStates <- c("sunny", "cloudy", "rain")
R> byRow <- TRUE
R> weatherMatrix <- matrix(data = c(0.7, 0.2, 0.1,
+                                   0.3, 0.4, 0.3,
+                                   0.2, 0.4, 0.4),
+                            byrow = byRow, nrow = 3,
+                            dimnames = list(weatherStates, weatherStates))
R> mcWeather <- new("markovchain", states = weatherStates,
+                 byrow = byRow, transitionMatrix = weatherMatrix,
+                 name = "Weather")
R> weathersOfDays <- rmarkovchain(n = 365, object = mcWeather, t0 = "sunny")
```

For the purpose of this section, we shall continue to use the weather of days example introduced in the main vignette of the package (reproduced above for convenience).

Let us invoke the fit function to estimate the MAP parameters with 92% confidence bounds and hyper-parameters as shown below, based on the first 200 days of the weather data. Additionally, let us find out what the probability is of observing the weather data for the next 165 days. The usage would be as follows

```
R> hyperMatrix<-matrix(c(1, 1, 2,
+                        3, 2, 1,
+                        2, 2, 3),
+                       nrow = 3, byrow = TRUE,
+                       dimnames = list(weatherStates,weatherStates))
R> markovchainFit(weathersOfDays[1:200], method = "map",
+               confidencelevel = 0.92, hyperparam = hyperMatrix)


$estimate
Bayesian Fit
 A  3 - dimensional discrete Markov Chain defined by the following states:
 cloudy, rain, sunny
 The transition matrix  (by rows)  is defined as follows:
          cloudy      rain      sunny
cloudy 0.3709677 0.2741935 0.3548387
rain   0.3809524 0.3809524 0.2380952
sunny  0.2135922 0.0776699 0.7087379



$expectedValue
          cloudy       rain      sunny
cloudy 0.3692308 0.27692308 0.3538462
rain   0.3777778 0.37777778 0.2444444
sunny  0.2169811 0.08490566 0.6981132


$standardError
            [,1]        [,2]        [,3]
```

```
[1,] 0.05940353 0.05508075 0.05885769
[2,] 0.07148453 0.07148453 0.06336426
[3,] 0.03984788 0.02694693 0.04438059


$confidenceInterval
$confidenceInterval$confidenceLevel
[1] 0.92


$confidenceInterval$lowerEndpointMatrix
          [,1]      [,2]      [,3]
[1,] 0.2788643 0.1837403 0.2628715
[2,] 0.2728115 0.2728115 0.1312487
[3,] 0.1436350 0.0000000 0.6344429


$confidenceInterval$upperEndpointMatrix
          [,1]      [,2]      [,3]
[1,] 0.4955338 0.3757588 0.4749592
[2,] 0.5431506 0.5431506 0.3509441
[3,] 0.2829567 0.1249744 1.0000000



$logLikelihood
[1] -182.4175


R> predictiveDistribution(weathersOfDays[1:200],
+                         weathersOfDays[201:365],hyperparam = hyperMatrix)


[1] -151.6332
```

The results should not change after permuting the dimensions of the matrix.

```
R> hyperMatrix2<- hyperMatrix[c(2,3,1), c(2,3,1)]
R> markovchainFit(weathersOfDays[1:200], method = "map",
+                 confidencelevel = 0.92, hyperparam = hyperMatrix2)


$estimate
Bayesian Fit
 A  3 - dimensional discrete Markov Chain defined by the following states:
 cloudy, rain, sunny
 The transition matrix  (by rows)  is defined as follows:
           cloudy      rain     sunny
cloudy 0.3709677 0.2741935 0.3548387
rain   0.3809524 0.3809524 0.2380952
sunny  0.2135922 0.0776699 0.7087379
```

```
$expectedValue
          cloudy      rain      sunny
cloudy 0.3692308 0.27692308 0.3538462
rain   0.3777778 0.37777778 0.2444444
sunny  0.2169811 0.08490566 0.6981132


$standardError
            [,1]        [,2]        [,3]
[1,] 0.05940353 0.05508075 0.05885769
[2,] 0.07148453 0.07148453 0.06336426
[3,] 0.03984788 0.02694693 0.04438059


$confidenceInterval
$confidenceInterval$confidenceLevel
[1] 0.92


$confidenceInterval$lowerEndpointMatrix
          [,1]      [,2]      [,3]
[1,] 0.2788643 0.1837403 0.2628715
[2,] 0.2728115 0.2728115 0.1312487
[3,] 0.1436350 0.0000000 0.6344429


$confidenceInterval$upperEndpointMatrix
          [,1]      [,2]      [,3]
[1,] 0.4955338 0.3757588 0.4749592
[2,] 0.5431506 0.5431506 0.3509441
[3,] 0.2829567 0.1249744 1.0000000



$logLikelihood
[1] -182.4175

R> predictiveDistribution(weathersOfDays[1:200],
+                   weathersOfDays[201:365],hyperparam = hyperMatrix2)

[1] -151.6332
```

Note that the predictive probability is very small. However, this can be useful when comparing model orders. Suppose we have an idea of the (prior) transition matrix corresponding to the expected value of the parameters, and have a data set from which we want to deduce the MAP estimates. We can infer the hyper-parameters from this known transition matrix itself, and use this to obtain our MAP estimates.

```
R> inferHyperparam(transMatr = weatherMatrix, scale = c(10, 10, 10))

$scaledInference
      cloudy rain sunny
```

```
cloudy     4    3    3
rain       4    4    2
sunny      2    1    7
```

Alternatively, we can use a data sample to infer the hyper-parameters.

```
R> inferHyperparam(data = weathersOfDays[1:15])
```

```
$dataInference
       cloudy rain sunny
cloudy      3    2     3
rain        3    4     1
sunny       2    1     4
```

In order to use the inferred hyper-parameter matrices, we do

```
R> hyperMatrix3 <- inferHyperparam(transMatr = weatherMatrix,
+                                  scale = c(10, 10, 10))
R> hyperMatrix3 <- hyperMatrix3$scaledInference
R> hyperMatrix4 <- inferHyperparam(data = weathersOfDays[1:15])
R> hyperMatrix4 <- hyperMatrix4$dataInference
```

Now we can safely use `hyperMatrix3` and `hyperMatrix4` with `markovchainFit` (in the `hyperparam` argument).

Supposing we don't provide any hyper-parameters, then the prior is uniform. This is the same as maximum likelihood.

```
R> data(preproglucacon)
R> preproglucacon <- preproglucacon[[2]]
R> MLEest <- markovchainFit(preproglucacon, method = "mle")
R> MAPest <- markovchainFit(preproglucacon, method = "map")
R> MLEest$estimate
```

```
MLE Fit
 A  4 - dimensional discrete Markov Chain defined by the following states:
 A, C, G, T
 The transition matrix  (by rows)  is defined as follows:
          A         C          G         T
A 0.3585271 0.1434109 0.16666667 0.3313953
C 0.3840304 0.1558935 0.02281369 0.4372624
G 0.3053097 0.1991150 0.15044248 0.3451327
T 0.2844523 0.1819788 0.17667845 0.3568905
```

```
R> MAPest$estimate
```

```
Bayesian Fit
 A  4 - dimensional discrete Markov Chain defined by the following states:
 A, C, G, T
 The transition matrix  (by rows)  is defined as follows:
            A           C           G           T
A 0.3585271 0.1434109 0.16666667 0.3313953
C 0.3840304 0.1558935 0.02281369 0.4372624
G 0.3053097 0.1991150 0.15044248 0.3451327
T 0.2844523 0.1819788 0.17667845 0.3568905
```

6. Applications

This section shows applications of DTMC in various fields.

6.1. Weather forecasting

Markov chains provide a simple model to predict the next day's weather given the current meteorological condition. The first application herewith shown is the "Land of Oz example" from J. G. Kemeny, J. L.Snell, and G. L. Thompson (1974), the second is the "Alofi Island Rainfall" from P. J. Avery and D. A. Henderson (1999).

Land of Oz

The Land of Oz is acknowledged not to have ideal weather conditions at all: the weather is snowy or rainy very often and, once more, there are never two nice days in a row. Consider three weather states: rainy, nice and snowy. Let the transition matrix be as in the following:

```
R> mcWP <- new("markovchain", states = c("rainy", "nice", "snowy"),
+          transitionMatrix = matrix(c(0.5, 0.25, 0.25,
+                                      0.5, 0, 0.5,
+                                      0.25,0.25,0.5), byrow = T, nrow = 3))
```

Given that today it is a nice day, the corresponding stochastic row vector is $w_0 = (0,1,0)$ and the forecast after 1, 2 and 3 days are given by

```
R> W0 <- t(as.matrix(c(0, 1, 0)))
R> W1 <- W0 * mcWP; W1


     rainy nice snowy
[1,]   0.5    0   0.5


R> W2 <- W0 * (mcWP ^ 2); W2


     rainy nice snowy
[1,] 0.375 0.25 0.375
```

```
R> W3 <- W0 * (mcWP ^ 3); W3
```

```
       rainy   nice   snowy
[1,] 0.40625 0.1875 0.40625
```

As can be seen from w_1, if in the Land of Oz today is a nice day, tomorrow it will rain or snow with probability 1. One week later, the prediction can be computed as

```
R> W7 <- W0 * (mcWP ^ 7)
R> W7
```

```
         rainy      nice     snowy
[1,] 0.4000244 0.1999512 0.4000244
```

The steady state of the chain can be computed by means of the `steadyStates` method.

```
R> q <- steadyStates(mcWP)
R> q
```

```
     rainy nice snowy
[1,]   0.4  0.2   0.4
```

Note that, from the seventh day on, the predicted probabilities are substantially equal to the steady state of the chain and they don't depend from the starting point, as the following code shows.

```
R> R0 <- t(as.matrix(c(1, 0, 0)))
R> R7 <- R0 * (mcWP ^ 7); R7
```

```
         rainy      nice     snowy
[1,] 0.4000244 0.2000122 0.3999634
```

```
R> S0 <- t(as.matrix(c(0, 0, 1)))
R> S7 <- S0 * (mcWP ^ 7); S7
```

```
         rainy      nice     snowy
[1,] 0.3999634 0.2000122 0.4000244
```

Alofi Island Rainfall

Alofi Island daily rainfall data were recorded from January 1st, 1987 until December 31st, 1989 and classified into three states: "0" (no rain), "1-5" (from non zero until 5 mm) and "6+" (more than 5mm). The corresponding dataset is provided within the **markovchain** package.

```
R> data("rain", package = "markovchain")
R> table(rain$rain)
```

```
   0 1-5  6+
548 295 253
```

The underlying transition matrix is estimated as follows.

```
R> mcAlofi <- markovchainFit(data = rain$rain, name = "Alofi MC")$estimate
R> mcAlofi

Alofi MC
 A  3 - dimensional discrete Markov Chain defined by the following states:
 0, 1-5, 6+
 The transition matrix  (by rows)  is defined as follows:
            0        1-5         6+
0   0.6605839 0.2299270 0.1094891
1-5 0.4625850 0.3061224 0.2312925
6+  0.1976285 0.3122530 0.4901186
```

The long term daily rainfall distribution is obtained by means of the `steadyStates` method.

```
R> steadyStates(mcAlofi)

              0        1-5         6+
[1,] 0.5008871 0.2693656 0.2297473
```

6.2. Finance and Economics

Other relevant applications of DTMC can be found in Finance and Economics.

Finance

Credit ratings transitions have been successfully modeled with discrete time Markov chains. Some rating agencies publish transition matrices that show the empirical transition probabilities across credit ratings. The example that follows comes from **CreditMetrics** R package (Wittmann 2007), carrying Standard & Poor's published data.

```
R> rc <- c("AAA", "AA", "A", "BBB", "BB", "B", "CCC", "D")
R> creditMatrix <- matrix(
+   c(90.81, 8.33, 0.68, 0.06, 0.08, 0.02, 0.01, 0.01,
+     0.70, 90.65, 7.79, 0.64, 0.06, 0.13, 0.02, 0.01,
+     0.09, 2.27, 91.05, 5.52, 0.74, 0.26, 0.01, 0.06,
+     0.02, 0.33, 5.95, 85.93, 5.30, 1.17, 1.12, 0.18,
+     0.03, 0.14, 0.67, 7.73, 80.53, 8.84, 1.00, 1.06,
+     0.01, 0.11, 0.24, 0.43, 6.48, 83.46, 4.07, 5.20,
+     0.21, 0, 0.22, 1.30, 2.38, 11.24, 64.86, 19.79,
+     0, 0, 0, 0, 0, 0, 0, 100
+     )/100, 8, 8, dimnames = list(rc, rc), byrow = TRUE)
```

It is easy to convert such matrices into `markovchain` objects and to perform some analyses

```
R> creditMc <- new("markovchain", transitionMatrix = creditMatrix,
+                  name = "S&P Matrix")
R> absorbingStates(creditMc)
```

```
[1] "D"
```

Economics

For a recent application of **markovchain** in Economic, see Jacob (2014).

A dynamic system generates two kinds of economic effects (Bard 2000):

1. those incurred when the system is in a specified state, and
2. those incurred when the system makes a transition from one state to another.

Let the monetary amount of being in a particular state be represented as a m-dimensional column vector c^S, while let the monetary amount of a transition be embodied in a C^R matrix in which each component specifies the monetary amount of going from state i to state j in a single step. Henceforth, Equation (24) represents the monetary of being in state i.

$$c_i = c_i^S + \sum_{j=1}^{m} C_{ij}^R p_{ij}. \tag{24}$$

Let $\bar{c} = [c_i]$ and let e_i be the vector valued 1 in the initial state and 0 in all other, then, if f_n is the random variable representing the economic return associated with the stochastic process at time n, Equation (25) holds:

$$E\left[f_n\left(X_n\right)|X_0 = i\right] = e_i P^n \bar{c}. \tag{25}$$

The following example assumes that a telephone company models the transition probabilities between customer/non-customer status by matrix P and the cost associated to states by matrix M.

```
R> statesNames <- c("customer", "non customer")
R> P <- zeros(2); P[1, 1] <- .9; P[1, 2] <- .1; P[2, 2] <- .95; P[2, 1] <- .05;
R> rownames(P) <- statesNames; colnames(P) <- statesNames
R> mcP <- new("markovchain", transitionMatrix = P, name = "Telephone company")
R> M <- zeros(2); M[1, 1] <- -20; M[1, 2] <- -30; M[2, 1] <- -40; M[2, 2] <- 0
```

If the average revenue for existing customer is +100, the cost per state is computed as follows.

```
R> c1 <- 100 + conditionalDistribution(mcP, state = "customer") %*% M[1,]
R> c2 <- 0 + conditionalDistribution(mcP, state = "non customer") %*% M[2,]
```

For an existing customer, the expected gain (loss) at the fifth year is given by the following code.

```
R> as.numeric((c(1, 0)* mcP ^ 5) %*% (as.vector(c(c1, c2))))
```

```
[1] 48.96009
```

6.3. Actuarial science

Markov chains are widely applied in the field of actuarial science. Two classical applications are policyholders' distribution across Bonus Malus classes in Motor Third Party Liability (MTPL) insurance (Section 6.3.1) and health insurance pricing and reserving (Section 6.3.2).

MPTL Bonus Malus

Bonus Malus (BM) contracts grant the policyholder a discount (enworsen) as a function of the number of claims in the experience period. The discount (enworsen) is applied on a premium that already allows for known (a priori) policyholder characteristics (Denuit, Maréchal, Pitrebois, and Walhin 2007) and it usually depends on vehicle, territory, the demographic profile of the policyholder, and policy coverage deep (deductible and policy limits).\ Since the proposed BM level depends on the claim on the previous period, it can be modeled by a discrete Markov chain. A very simplified example follows. Assume a BM scale from 1 to 5, where 4 is the starting level. The evolution rules are shown in Equation 26:

$$bm_{t+1} = \max\left(1, bm_t - 1\right) * \left(\tilde{N} = 0\right) + \min\left(5, bm_t + 2 * \tilde{N}\right) * \left(\tilde{N} \geq 1\right). \tag{26}$$

The number of claim $\tilde{N}$ is a random variable that is assumed to be Poisson distributed.

```
R> getBonusMalusMarkovChain <- function(lambda) {
+    bmMatr <- zeros(5)
+    bmMatr[1, 1] <- dpois(x = 0, lambda)
+    bmMatr[1, 3] <- dpois(x = 1, lambda)
+    bmMatr[1, 5] <- 1 - ppois(q = 1, lambda)
+
+    bmMatr[2, 1] <- dpois(x = 0, lambda)
+    bmMatr[2, 4] <- dpois(x = 1, lambda)
+    bmMatr[2, 5] <- 1 - ppois(q = 1, lambda)
+
+    bmMatr[3, 2] <- dpois(x = 0, lambda)
+    bmMatr[3, 5] <- 1 - dpois(x=0, lambda)
+
+    bmMatr[4, 3] <- dpois(x = 0, lambda)
+    bmMatr[4, 5] <- 1 - dpois(x = 0, lambda)
+
+    bmMatr[5, 4] <- dpois(x = 0, lambda)
+    bmMatr[5, 5] <- 1 - dpois(x = 0, lambda)
```

```
+   stateNames <- as.character(1:5)
+   out <- new("markovchain", transitionMatrix = bmMatr,
+             states = stateNames, name = "BM Matrix")
+   return(out)
+ }
```

Assuming that the a-priori claim frequency per car-year is 0.05 in the class (being the class the group of policyholders that share the same common characteristics), the underlying BM transition matrix and its underlying steady state are as follows.

```
R> bmMc <- getBonusMalusMarkovChain(0.05)
R> as.numeric(steadyStates(bmMc))
```

```
[1] 0.895836079 0.045930498 0.048285405 0.005969247 0.003978772
```

If the underlying BM coefficients of the class are 0.5, 0.7, 0.9, 1.0, 1.25, this means that the average BM coefficient applied on the long run to the class is given by

```
R> sum(as.numeric(steadyStates(bmMc)) * c(0.5, 0.7, 0.9, 1, 1.25))
```

```
[1] 0.534469
```

This means that the average premium paid by policyholders in the portfolio almost halves in the long run.

Health insurance example

Actuaries quantify the risk inherent in insurance contracts evaluating the premium of insurance contract to be sold (therefore covering future risk) and evaluating the actuarial reserves of existing portfolios (the liabilities in terms of benefits or claims payments due to policyholder arising from previously sold contracts), see Deshmukh (2012) for details.

An applied example can be performed using the data from De Angelis, Paolo and Di Falco, L. (2016) that has been saved in the `exdata` folder.

```
R> ltcDemoPath<-system.file("extdata", "ltdItaData.txt",
+                        package = "markovchain")
R> ltcDemo<-read.table(file = ltcDemoPath, header=TRUE,
+                   sep = ";", dec = ".")
R> head(ltcDemo)
```

```
  age         pAD        pID         pAI        pAA
1  20 0.0004616002 0.01083364 0.0001762467 0.9993622
2  21 0.0004824888 0.01079719 0.0001710577 0.9993465
3  22 0.0004949938 0.01177076 0.0001592333 0.9993458
4  23 0.0005042935 0.01159394 0.0001605731 0.9993351
5  24 0.0005074193 0.01260574 0.0001606504 0.9993319
6  25 0.0005154267 0.01526364 0.0001643603 0.9993202
```

The data shows the probability of transition between the state of (A)ctive, to (I)ll and Dead. It is easy to complete the transition matrix.

```
R> ltcDemo<-transform(ltcDemo,
+                   pIA=0,
+                   pII=1-pID,
+                   pDD=1,
+                   pDA=0,
+                   pDI=0)
```

Now we build a function that returns the transition during the $t+1$ th year, assuming that the subject has attained year t.

```
R> possibleStates<-c("A","I","D")
R> getMc4Age<-function(age) {
+    transitionsAtAge<-ltcDemo[ltcDemo$age==age,]
+
+    myTransMatr<-matrix(0, nrow=3,ncol = 3,
+                      dimnames = list(possibleStates, possibleStates))
+    myTransMatr[1,1]<-transitionsAtAge$pAA[1]
+    myTransMatr[1,2]<-transitionsAtAge$pAI[1]
+    myTransMatr[1,3]<-transitionsAtAge$pAD[1]
+    myTransMatr[2,2]<-transitionsAtAge$pII[1]
+    myTransMatr[2,3]<-transitionsAtAge$pID[1]
+    myTransMatr[3,3]<-1
+
+    myMc<-new("markovchain", transitionMatrix = myTransMatr,
+            states = possibleStates,
+            name = paste("Age",age,"transition matrix"))
+
+    return(myMc)
+
+ }
```

Cause transitions are not homogeneous across ages, we use a `markovchainList` object to describe the transition probabilities for a guy starting at age 100.

```
R> getFullTransitionTable<-function(age){
+    ageSequence<-seq(from=age, to=120)
+    k=1
+    myList=list()
+    for ( i in ageSequence) {
+      mc_age_i<-getMc4Age(age = i)
+      myList[[k]]<-mc_age_i
+      k=k+1
+    }
```

```
+    myMarkovChainList<-new("markovchainList", markovchains = myList,
+                          name = paste("TransitionsSinceAge", age, sep = ""))
+    return(myMarkovChainList)
+ }
R> transitionsSince100<-getFullTransitionTable(age=100)
```

We can use such transition for simulating ten life trajectories for a guy that begins "active" (A) aged 100:

```
R> rmarkovchain(n = 10, object = transitionsSince100,
+              what = "matrix", t0 = "A", include.t0 = TRUE)
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13]
 [1,] "A"  "A"  "A"  "D"  "D"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
 [2,] "A"  "A"  "A"  "D"  "D"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
 [3,] "A"  "A"  "D"  "D"  "D"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
 [4,] "A"  "A"  "I"  "I"  "D"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
 [5,] "A"  "I"  "D"  "D"  "D"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
 [6,] "A"  "A"  "A"  "A"  "A"  "A"  "A"  "A"  "A"  "I"   "D"   "D"   "D"
 [7,] "A"  "A"  "A"  "A"  "A"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
 [8,] "A"  "A"  "A"  "I"  "D"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
 [9,] "A"  "A"  "A"  "A"  "D"  "D"  "D"  "D"  "D"  "D"   "D"   "D"   "D"
[10,] "A"  "A"  "A"  "A"  "A"  "A"  "I"  "D"  "D"  "D"   "D"   "D"   "D"
      [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22]
 [1,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [2,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [3,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [4,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [5,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [6,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [7,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [8,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
 [9,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
[10,] "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"   "D"
```

Lets consider 1000 simulated live trajectories, for a healthy guy aged 80. We can compute the expected time a guy will be disabled starting active at age 80.

```
R> transitionsSince80<-getFullTransitionTable(age=80)
R> lifeTrajectories<-rmarkovchain(n=1e3, object=transitionsSince80,
+                                 what="matrix",t0="A",include.t0=TRUE)
R> temp<-matrix(0,nrow=nrow(lifeTrajectories),ncol = ncol(lifeTrajectories))
R> temp[lifeTrajectories=="I"]<-1
R> expected_period_disabled<-mean(rowSums((temp)))
R> expected_period_disabled
```

```
[1] 1.361
```

Assuming that the health insurance will pay a benefit of 12000 per year disabled and that the real interest rate is 0.02, we can compute the lump sum premium at 80.

```
R> mean(rowMeans(12000*temp%*%( matrix((1+0.02)^-seq(from=0, to=ncol(temp)-1)))))

[1] 13621.67
```

6.4. Sociology

Markov chains have been actively used to model progressions and regressions between social classes. The first study was performed by Glass and Hall (1954), while a more recent application can be found in Jo Blanden and Machin (2005). The table that follows shows the income quartile of the father when the son was 16 (in 1984) and the income quartile of the son when aged 30 (in 2000) for the 1970 cohort.

```
R> data("blanden")
R> mobilityMc <- as(blanden, "markovchain")
R> mobilityMc

Unnamed Markov chain
 A  4 - dimensional discrete Markov Chain defined by the following states:
 Bottom, 2nd, 3rd, Top
 The transition matrix  (by rows)  is defined as follows:
             2nd       3rd      Bottom        Top
Bottom 0.2900000 0.2200000 0.3800000 0.1100000
2nd     0.2772277 0.2574257 0.2475248 0.2178218
3rd     0.2626263 0.2828283 0.2121212 0.2424242
Top     0.1700000 0.2500000 0.1600000 0.4200000
```

The underlying transition graph is plotted in Figure 5.

The steady state distribution is computed as follows. Since transition across quartiles are shown, the probability function is evenly 0.25.

```
R> round(steadyStates(mobilityMc), 2)

     Bottom  2nd  3rd  Top
[1,]   0.25 0.25 0.25 0.25
```

6.5. Genetics and Medicine

This section contains two examples: the first shows the use of Markov chain models in genetics, the second shows an application of Markov chains in modelling diseases' dynamics.

Genetics

P. J. Avery and D. A. Henderson (1999) discusses the use of Markov chains in model Preproglucacon gene protein bases sequence. The `preproglucacon` dataset in **markovchain** contains the dataset shown in the package.

1970 mobility

Figure 5: 1970 UK cohort mobility data.

```
R> data("preproglucacon", package = "markovchain")
```

It is possible to model the transition probabilities between bases as shown in the following code.

```
R> mcProtein <- markovchainFit(preproglucacon$preproglucacon,
+                          name = "Preproglucacon MC")$estimate
R> mcProtein

Preproglucacon MC
 A  4 - dimensional discrete Markov Chain defined by the following states:
 A, C, G, T
 The transition matrix  (by rows)  is defined as follows:
          A         C          G         T
A 0.3585271 0.1434109 0.16666667 0.3313953
C 0.3840304 0.1558935 0.02281369 0.4372624
G 0.3053097 0.1991150 0.15044248 0.3451327
T 0.2844523 0.1819788 0.17667845 0.3568905
```

Medicine

Discrete-time Markov chains are also employed to study the progression of chronic diseases. The following example is taken from B. A. Craig and A. A. Sendi (2002). Starting from six

month follow-up data, the maximum likelihood estimation of the monthly transition matrix is obtained. This transition matrix aims to describe the monthly progression of CD4-cell counts of HIV infected subjects.

```
R> craigSendiMatr <- matrix(c(682, 33, 25,
+               154, 64, 47,
+               19, 19, 43), byrow = T, nrow = 3)
R> hivStates <- c("0-49", "50-74", "75-UP")
R> rownames(craigSendiMatr) <- hivStates
R> colnames(craigSendiMatr) <- hivStates
R> craigSendiTable <- as.table(craigSendiMatr)
R> mcM6 <- as(craigSendiTable, "markovchain")
R> mcM6@name <- "Zero-Six month CD4 cells transition"
R> mcM6


Zero-Six month CD4 cells transition
 A  3 - dimensional discrete Markov Chain defined by the following states:
 0-49, 50-74, 75-UP
 The transition matrix  (by rows)  is defined as follows:
           0-49        50-74        75-UP
0-49   0.9216216 0.04459459 0.03378378
50-74 0.5811321 0.24150943 0.17735849
75-UP 0.2345679 0.23456790 0.53086420
```

As shown in the paper, the second passage consists in the decomposition of $M_6 = V \cdot D \cdot V^{-1}$ in order to obtain M_1 as $M_1 = V \cdot D^{1/6} \cdot V^{-1}$.

```
R> eig <- eigen(mcM6@transitionMatrix)
R> D <- diag(eig$values)

R> V <- eig$vectors
R> V %*% D %*% solve(V)


          [,1]        [,2]        [,3]
[1,] 0.9216216 0.04459459 0.03378378
[2,] 0.5811321 0.24150943 0.17735849
[3,] 0.2345679 0.23456790 0.53086420


R> d <- D ^ (1/6)
R> M <- V %*% d %*% solve(V)
R> mcM1 <- new("markovchain", transitionMatrix = M, states = hivStates)
```

7. Discussion, issues and future plans

The **markovchain** package has been designed in order to provide easily handling of DTMC and communication with alternative packages.

The package has known several improvements in the recent years: many functions added, porting the software in Rcpp **Rcpp** package (Eddelbuettel 2013) and many methodological improvements that have improved the software reliability.

8. Acknowledgments

References

Allaire J, Francois R, Ushey K, Vandenbrouck G, Geelnard M, Intel (2016). *RcppParallel: Parallel Programming Tools for 'Rcpp'*. R package version 4.3.19, URL http://rcppcore.github.io/RcppParallel/.

Anderson TW, Goodman LA (1957). "Statistical inference about Markov chains." *The Annals of Mathematical Statistics*, pp. 89–110.

B A Craig, A A Sendi (2002). "Estimation of the Transition Matrix of a Discrete-Time Markov Chain." *Health Economics*, **11**, 33–42.

Bard JF (2000). "Lecture 12.5 - Additional Issues Concerning Discrete-Time Markov Chains." URL http://www.me.utexas.edu/~jensen%20/ORMM/instruction/powerpoint/or_models_09/12.5_dtmc2.ppt.

Brémaud P (1999). "Discrete-Time Markov Models." In *Markov Chains*, pp. 53–93. Springer.

Chambers J (2008). *Software for Data Analysis: Programming with R*. Statistics and computing. Springer-Verlag. ISBN 9780387759357.

Ching WK, Ng MK, Fung ES (2008). "Higher-order multivariate Markov chains and their applications." *Linear Algebra and its Applications*, **428**(2), 492–507.

Csardi G, Nepusz T (2006). "The **igraph** Software Package for Complex Network Research." *InterJournal*, **Complex Systems**, 1695. URL http://igraph.sf.net.

De Angelis, Paolo, Di Falco, L (2016). *Assicurazioni sulla salute: caratteristiche, modelli attuariali e basi tecniche*. Il Mulino. Il Mulino. ISBN 9788815260840. URL https://books.google.it/books?id=D56bjgEACAAJ.

de Wreede LC, Fiocco M, Putter H (2011). "**mstate**: An R Package for the Analysis of Competing Risks and Multi-State Models." *Journal of Statistical Software*, **38**(7), 1–30. URL http://www.jstatsoft.org/v38/i07/.

Denuit M, Maréchal X, Pitrebois S, Walhin JF (2007). *Actuarial modelling of claim counts: Risk classification, credibility and bonus-malus systems*. Wiley.

Deshmukh S (2012). *Multiple Decrement Models in Insurance: An Introduction Using R*. SpringerLink : Bücher. Springer-Verlag. ISBN 9788132206590.

Dobrow RP (2016). *Introduction to Stochastic Processes with R*. John Wiley & Sons.

Eddelbuettel D (2013). *Seamless R and C++ Integration with **Rcpp***. Springer-Verlag, New York. ISBN 978-1-4614-6867-7.

Feres R (2007). "Notes for Math 450 MATLAB Listings for Markov Chains." URL http://www.math.wustl.edu/~feres/Math450Lect04.pdf.

Geyer CJ, Johnson LT (2013). *mcmc: Markov Chain Monte Carlo*. R package version 0.9-2, URL http://CRAN.R-project.org/package=mcmc.

Glass D, Hall JR (1954). "Social Mobility in Great Britain: A Study in Intergenerational Change in Status." In *Social Mobility in Great Britain*. Routledge and Kegan Paul.

Goulet V, Dutang C, Maechler M, Firth D, Shapira M, Stadelmann M, expm-developers@listsR-forgeR-projectorg (2013). **expm**: *Matrix Exponential*. R package version 0.99-1, URL http://CRAN.R-project.org/package=expm.

Grinstead CM, Snell LJ (2006). *Grinstead and Snell's Introduction to Probability*. Version dated 4 july 2006 edition. American Mathematical Society. URL http://math.dartmouth.edu/~{}prob/prob/prob.pdf.

Himmelmann SSDDL, wwwlinhicom (2010). **HMM**: *HMM - Hidden Markov Models*. R package version 1.0, URL http://CRAN.R-project.org/package=HMM.

Hu YT, Kiesel R, Perraudin W (2002). "The estimation of transition matrices for sovereign credit ratings." *Journal of Banking and Finance*, **26**(7), 1383–1406. ISSN 03784266. doi:10.1016/S0378-4266(02)00268-6.

Israel RB, Rosenthal JS, Wei JZ (2001). "Finding generators for Markov chains via empirical transition matrices, with applications to credit ratings." *Mathematical finance*, **11**(2), 245–265.

J G Kemeny, J LSnell, G L Thompson (1974). *Introduction to Finite Mathematics*. Prentice Hall.

Jackson CH (2011). "Multi-State Models for Panel Data: The **msm** Package for R." *Journal of Statistical Software*, **38**(8), 1–29. URL http://www.jstatsoft.org/v38/i08/.

Jacob I (2014). "Is R Cost Effective?" Electronic. Presented on Manchester R Meeting, URL http://www.rmanchester.org/Presentations/Ian%20Jacob%20-%20Is%20R%20Cost%20Effective.pdf.

Jo Blanden PG, Machin S (2005). "Intergenerational Mobility in Europe and North America." *Technical report*, Center for Economic Performances. URL http://cep.lse.ac.uk/about/news/IntergenerationalMobility.pdf.

Konstantopoulos T (2009). "Markov Chains and Random Walks." *Lecture notes*.

Kullback S, Kupperman M, Ku H (1962). "Tests for Contingency Tables and Marltov Chains." *Technometrics*, **4**(4), 573–608.

Montgomery J (2009). "Communication Classes." URL http://www.ssc.wisc.edu/~jmontgom/commclasses.pdf.

Nicholson W (2013). **DTMCPack**: *Suite of Functions Related to Discrete-Time Discrete-State Markov Chains*. R package version 0.1-2, URL http://CRAN.R-project.org/package=DTMCPack.

Noé F, Schütte C, Vanden-Eijnden E, Reich L, Weikl TR (2009). "Constructing the equilibrium ensemble of folding pathways from short off-equilibrium simulations." *Proceedings of the National Academy of Sciences*, **106**(45), 19011–19016. ISSN 0027-8424, 1091-6490. doi:10.1073/pnas.0905466106. URL https://www.pnas.org/content/106/45/19011.

P J Avery, D A Henderson (1999). "Fitting Markov Chain Models to Discrete State Series." *Applied Statistics*, **48**(1), 53–61.

Pfeuffer M (2017). *ctmcd: Estimating the Parameters of a Continuous-Time Markov Chain from Discrete-Time Data.* R package version 1.1, URL https://CRAN.R-project.org/package=ctmcd.

R Core Team (2013). *R: A Language and Environment for Statistical Computing.* R Foundation for Statistical Computing, Vienna, Austria. URL http://www.R-project.org/.

Roebuck P (2011). **matlab***: MATLAB emulation package.* R package version 0.8.9, URL http://CRAN.R-project.org/package=matlab.

Sison CP, Glaz J (1995). "Simultaneous confidence intervals and sample size determination for multinomial proportions." *Journal of the American Statistical Association*, **90**(429), 366–369.

Skuriat-Olechnowska M (2005). *Statistical inference and hypothesis testing for Markov chains with Interval Censoring.* diploma thesis, Delft University of Technology.

Snell L (1999). "Probability Book: chapter 11." URL http://www.dartmouth.edu/~chance/teaching_aids/books_articles/probability_book/Chapter11.pdf.

Spedicato GA (2017). "Discrete Time Markov Chains with R." *The R Journal*. URL https://journal.r-project.org/archive/2017/RJ-2017-036/index.html.

Tarjan R (1972). "Depth-First Search and Linear Graph Algorithms." *SIAM Journal on Computing*, **1**(2), 146–160. doi:10.1137/0201010. URL https://doi.org/10.1137/0201010.

Villacorta PJ (2012). *MultinomialCI: Simultaneous confidence intervals for multinomial proportions according to the method by Sison and Glaz.* R package version 1.0, URL http://CRAN.R-project.org/package=MultinomialCI.

Wittmann A (2007). *CreditMetrics: Functions for Calculating the CreditMetrics Risk Model.* R package version 0.0-2.

Wolfram Research I (2013a). URL http://www.wolfram.com/mathematica/new-in-9/markov-chains-and-queues/structural-properties-of-finite-markov-processes.html.

Wolfram Research I (2013b). *Mathematica.* Wolfram Research, Inc., ninth edition.

Affiliation:

Giorgio Alfredo Spedicato
Ph.D C.Stat FCAS, FSA, CSPA Unipol Group
Via Firenze 11 Paderno Dugnano 20037 Italy
E-mail: spedygiorgio@gmail.com
URL: www.statisticaladvisor.com

Tae Seung Kang
Ph.D student, Computer & Information Science & Engineering
University of Florida Gainesville, FL, USA
E-mail: tskang3@gmail.com

Sai Bhargav Yalamanchi
B-Tech student, Electrical Engineering
Indian Institute of Technology, Bombay Mumbai - 400 076, India
E-mail: bhargavcoolboy@gmail.com

Deepak Yadav
B-Tech student, Computer Science and Engineering
Indian Institute of Technology, Varanasi Uttar Pradesh - 221 005, India
E-mail: deepakyadav.iitbhu@gmail.com

Ignacio Cordón
Software Engineer
Madrid (Madrid), Spain
E-mail: nacho.cordon.castillo@gmail.com